

HPC LAB - Matrix Addition MPI

Name: Paleti Krishnasai

Roll No: CED18I039

Programming Environment: MPI

Problem: Matrix Addition

Date: 21-10-2021

Hardware Configuration:

PU NAME: Intel(R) Core(TM) i5-8300H CPU @ 2.30GHz

Number of Sockets: 1

Cores per Socket: 4

Threads per core: 2

L1d cache: 128 KiB

L1i cache: 128 KiB

L2 cache: 1 MiB

L3 cache: 8 MiB

```

root@galetti-Lenovo-ideapad-350-15IGU: ~$ lshw -l topology
CPU name:      Intel(R) Core(TM) i5-8300H CPU @ 2.30GHz
CPU type:      Intel Coffeelake processor
CPU stepping:  10
=====
Hardware Thread Topology
=====
Sockets:      1
Cores per socket: 4
Threads per core: 2
=====
HMThread      Thread      Core      Socket      Available
0              0              0          0          *
1              0              1          0          *
2              0              2          0          *
3              0              3          0          *
4              1              0          0          *
5              1              1          0          *
6              1              2          0          *
7              1              3          0          *
=====
Socket 0:      ( 0 4 1 5 2 6 3 7 )
=====
Cache Topology
=====
Level:         1
Size:          32 kB
Cache groups:  ( 0 4 ) ( 1 5 ) ( 2 6 ) ( 3 7 )
Level:         2
Size:          256 kB
Cache groups:  ( 0 4 ) ( 1 5 ) ( 2 6 ) ( 3 7 )
Level:         3
Size:          8 MB
Cache groups:  ( 0 4 1 5 2 6 3 7 )
=====
NUMA Topology
=====
NUMA domains:  1
Domain:        0
Processors:    ( 0 1 2 3 4 5 6 7 )
Distances:     10
Free memory:   3546.2 MB
Total memory:  7831.84 MB

```

```

*****
Graphical Topology
*****
Socket 0:
+-----+
| +-----+ | +-----+ | +-----+ | +-----+ | | | | | | | | |
| | 0 4 | | | 1 5 | | | 2 6 | | | 3 7 | |
| +-----+ | +-----+ | +-----+ | +-----+ |
| +-----+ | +-----+ | +-----+ | +-----+ |
| | 32 kB | | | 32 kB | | | 32 kB | | | 32 kB | |
| +-----+ | +-----+ | +-----+ | +-----+ |
| +-----+ | +-----+ | +-----+ | +-----+ |
| | 256 kB | | | 256 kB | | | 256 kB | | | 256 kB | |
| +-----+ | +-----+ | +-----+ | +-----+ |
| | 8 MB | | | | | | | |
| +-----+ | +-----+ | +-----+ | +-----+ |
+-----+

```

No of nodes: 12 (4 for each as written in machine file).

Serial Code:

```
#include "mpi.h"

#include <stdio.h>

#include <stdlib.h>


#define ROWS 20                /* number of rows */

#define COLUMNS 20            /* number of columns */

#define MASTER 0               /* taskid of first task */

#define FROM_MASTER 1          /* setting a message type */

#define FROM_WORKER 2          /* setting a message type */


int main(int argc, char *argv[])

{

    double start,end;

    MPI_Init(&argc,&argv);

    start=MPI_Wtime();

    int i,j,k;

    long double a[ROWS][COLUMNS],          /* matrix A to be Added */

    b[COLUMNS][COLUMNS],                  /* matrix B to be Added */

    c[ROWS][COLUMNS];

    //printf("Initializing arrays...\n");
```

```
for (i=0; i<ROWS; i++)

    for (j=0; j<COLUMNS; j++)

        a[i][j]= (i+j)*1.6584;

for (i=0; i<COLUMNS; i++)

    for (j=0; j<COLUMNS; j++)

        b[i][j]= (i+j)*4.2367;

for (k=0; k<COLUMNS; k++)                // Matrix Addition

    for (i=0; i<ROWS; i++)

    {

        c[i][k] = a[i][k] + b[i][k];

    }

for (i=0; i<ROWS; i++)

    {

        printf("\n");

        for (j=0; j<COLUMNS; j++)

            printf("%6.2Lf    ", c[i][j]);}

printf("\nFinished.\n");

end=MPI_Wtime();

printf("\nTime= %f",end-start);

return 0;

}
```

Parallel Code :

```
#include "mpi.h"

#include <stdio.h>

#include <stdlib.h>


#define ROWS 200                /* number of rows */

#define COLUMNS 200            /* number of columns */

#define MASTER 0                /* taskid of first task */

#define FROM_MASTER 1           /* setting a message type */

#define FROM_WORKER 2           /* setting a message type */


int main (int argc, char *argv[])

{

int    numtasks,                /* number of tasks in partition */

    taskid,                    /* a task identifier */

    numworkers,                /* number of worker tasks */

    source,                    /* task id of message source */

    dest,                      /* task id of message destination */

    mtype,                     /* message type */

    rows,                      /* rows of matrix A sent to each worker */

    averow, extra, offset, /* used to determine rows sent to each worker */
```

```
i, j, k, rc;          /* misc */

long double a[ROWS][COLUMNS],          /* matrix A to be Added */

b[COLUMNS][COLUMNS],          /* matrix B to be Added */

c[ROWS][COLUMNS];          /* result matrix C */

MPI_Status status;

double start,end;


MPI_Init(&argc,&argv);

start=MPI_Wtime();

MPI_Comm_rank(MPI_COMM_WORLD,&taskid);

MPI_Comm_size(MPI_COMM_WORLD,&numtasks);

if (numtasks < 2 ) {

    printf("Need at least two MPI tasks. Quitting...\n");

    MPI_Abort(MPI_COMM_WORLD, rc);

    exit(1);

}

numworkers = numtasks-1;
```

```
// ----- master task
// -----//

if (taskid == MASTER)

{

    //printf("mpi_Matrix Addition has started with %d
tasks.\n", numtasks);

    //printf("Initializing arrays...\n");

    for (i=0; i<ROWS; i++)

        for (j=0; j<COLUMNS; j++)

            a[i][j] = (i+j)*1.6584;

    for (i=0; i<COLUMNS; i++)

        for (j=0; j<COLUMNS; j++)

            b[i][j] = (i+j)*4.2367;

    /* Send matrix data to the worker tasks */

    averow = ROWS/numworkers;

    extra = ROWS%numworkers;

    offset = 0;

    mtype = FROM_MASTER;

    for (dest=1; dest<=numworkers; dest++)
```

```
{

    rows = (dest <= extra) ? averow+1 : averow;

    //printf("Sending %d rows to task %d\n", rows, dest, offset);

    MPI_Send(&offset, 1, MPI_INT, dest, mtype, MPI_COMM_WORLD);

    MPI_Send(&rows, 1, MPI_INT, dest, mtype, MPI_COMM_WORLD);

    MPI_Send(&a[offset][0], rows*COLUMNNS, MPI_LONG_DOUBLE, dest,
mtype,MPI_COMM_WORLD);

    MPI_Send(&b[offset][0], rows*COLUMNNS, MPI_LONG_DOUBLE, dest,
mtype, MPI_COMM_WORLD);

    offset = offset + rows;

}

// ----- Receive results from worker tasks
-----

mtype = FROM_WORKER;

for (i=1; i<=numworkers; i++)

{

    source = i;

    MPI_Recv(&offset, 1, MPI_INT, source, mtype, MPI_COMM_WORLD,
&status);
```

```
        MPI_Recv(&rows, 1, MPI_INT, source, mtype, MPI_COMM_WORLD,
&status);

        MPI_Recv(&c[offset][0], rows*COLUMNS, MPI_LONG_DOUBLE, source,
mtype,

                MPI_COMM_WORLD, &status);

        printf("Received results from task %d\n",source);

    }

    // Print results

printf("-----\n");

    printf("Result Matrix:\n");

    for (i=0; i<ROWS; i++)

    {

        printf("\n");

        for (j=0; j<COLUMNS; j++)

            printf("%6.2Lf    ", c[i][j]);

    }

printf("\n-----\n");

    printf ("Done.\n");
```



```
        end=MPI_Wtime();

        printf("\nTime= %f",end-start);

    }

// ----- worker task
// -----

    if (taskid > MASTER)

    {

        mtype = FROM_MASTER;

        MPI_Recv(&offset, 1, MPI_INT, MASTER, mtype, MPI_COMM_WORLD,
&status);

        MPI_Recv(&rows, 1, MPI_INT, MASTER, mtype, MPI_COMM_WORLD, &status);

        MPI_Recv(&a, rows*COLUMNS, MPI_LONG_DOUBLE, MASTER, mtype,
MPI_COMM_WORLD, &status);

        MPI_Recv(&b, rows*COLUMNS, MPI_LONG_DOUBLE, MASTER, mtype,
MPI_COMM_WORLD, &status);

        for (k=0; k<COLUMNS; k++)            // Matrix Addition

            for (i=0; i<rows; i++)

            {
```

```
        c[i][k] = a[i][k] + b[i][k];

    }

    mtype = FROM_WORKER;

    MPI_Send(&offset, 1, MPI_INT, MASTER, mtype, MPI_COMM_WORLD);

    MPI_Send(&rows, 1, MPI_INT, MASTER, mtype, MPI_COMM_WORLD);

    MPI_Send(&c, rows*COLUMNS, MPI_LONG_DOUBLE, MASTER, mtype,
MPI_COMM_WORLD);

}


MPI_Finalize();

}
```


Observations:

Processes(n)	Runtime	Speedup (s)
1	0.051765	1
2	0.453605	0.1141191124
4	0.619022	0.08362384536
6	0.619272	0.08359008642
8	0.613912	0.08431990253
10	0.632624	0.08182585548
12	0.514897	0.10053467
16	0.949612	0.05451173743
20	1.112871	0.04651482517
32	1.909993	0.02710219357
64	6.402997	0.008084495432
128	19.987483	0.002589870871

Runtime vs. Processes(n)



Speedup (s) vs. Processes(n)



Inference: (**Note:** Execution time, graph, and inference will be based on hardware configuration)

- Since MPI is a distributed memory architecture, the communication overhead between nodes causes the parallel code to run slower compared to serial code (running in 1 node or only in master)