

HPC LAB - Matrix Multiplication MPI

Name: Paleti Krishnasai

Roll No: CED18I039

Programming Environment: MPI

Problem: Matrix Multiplication

Date: 21-10-2021

Hardware Configuration:

PU NAME: Intel(R) Core(TM) i5-8300H CPU @ 2.30GHz

Number of Sockets: 1

Cores per Socket: 4

Threads per core: 2

L1d cache: 128 KiB

L1i cache: 128 KiB

L2 cache: 1 MiB

L3 cache: 8 MiB

```
paleti@paleti-Lenovo-Ideapad-330-151CH:~$ lllwid-topology
-----
CPU name:      Intel(R) Core(TM) i5-8300H CPU @ 2.30GHz
CPU type:      Intel Coffeelake processor
CPU stepping:  10
-----
Hardware Thread Topology
-----
Sockets:      1
Cores per socket: 4
Threads per core: 2
-----
HWThread      Thread      Core      Socket      Available
-----
0              0              0          0            *
1              0              1          0            *
2              0              2          0            *
3              0              3          0            *
4              1              0          0            *
5              1              1          0            *
6              1              2          0            *
7              1              3          0            *
-----
Socket 0:      ( 0 4 1 5 2 6 3 7 )
-----
Cache Topology
-----
Level:         1
Size:          32 kB
Cache groups:  ( 0 4 ) ( 1 5 ) ( 2 6 ) ( 3 7 )
-----
Level:         2
Size:          256 kB
Cache groups:  ( 0 4 ) ( 1 5 ) ( 2 6 ) ( 3 7 )
-----
Level:         3
Size:          8 MB
Cache groups:  ( 0 4 1 5 2 6 3 7 )
-----
NUMA Topology
-----
NUMA domains:  1
-----
Domain:        0
Processors:    ( 0 1 2 3 4 5 6 7 )
Distances:     10
Free memory:   3546.2 MB
Total memory:  7831.84 MB
-----
```

```
*****
Graphical Topology
*****
Socket 0:
+-----+
| +-----+ +-----+ +-----+ +-----+ |
| | 0 4 | | 1 5 | | 2 6 | | 3 7 | |
| +-----+ +-----+ +-----+ +-----+ |
| | 32 kB | | 32 kB | | 32 kB | | 32 kB | |
| +-----+ +-----+ +-----+ +-----+ |
| | 256 kB | | 256 kB | | 256 kB | | 256 kB | |
| +-----+ +-----+ +-----+ +-----+ |
| |                                     | 8 MB | |
| +-----+ +-----+ +-----+ +-----+ |
+-----+
```

No of nodes: 12 (4 for each as written in machine file).

Serial Code:

```
#include "mpi.h"

#include <stdio.h>

#include <stdlib.h>


#define ROWS_A 500                /* number of rows in matrix A */

#define COLUMNS_A 300            /* number of columns in matrix A */

#define COLUMNS_B 150           /* number of columns in matrix B */

#define MASTER 0                  /* taskid of first task */

#define FROM_MASTER 1             /* setting a message type */

#define FROM_WORKER 2            /* setting a message type */

int main(int argc, char *argv[])

{

    double start,end;
MPI_Init(&argc,&argv);

    start=MPI_Wtime();

    int i,j,k;

    long double    a[ROWS_A][COLUMNS_A],          /* matrix A to be
multiplied */
    b[COLUMNS_A][COLUMNS_B],                      /* matrix B to be multiplied */
    c[ROWS_A][COLUMNS_B];

    //printf("Initializing arrays...\n");

    for (i=0; i<ROWS_A; i++)
```

```
        for (j=0; j<COLUMNS_A; j++)

            a[i][j]= i+j;

    for (i=0; i<COLUMNS_A; i++)

        for (j=0; j<COLUMNS_B; j++)

            b[i][j]= i*j;

    for (k=0; k<COLUMNS_B; k++)

        for (i=0; i<ROWS_A; i++)

        {

            c[i][k] = 0.0;

            for (j=0; j<COLUMNS_A; j++)

                c[i][k] = c[i][k] + a[i][j] * b[j][k];

        }

    printf("-----\n");
    printf("Result Matrix:\n");
    for (i=0; i<ROWS_A; i++)
    {
        printf("\n");

        for (j=0; j<COLUMNS_B; j++)
            printf("%6.2Lf    ", c[i][j]);

    }
    printf("\n-----\n");
    printf ("Done.\n");
    end=MPI_Wtime();
    printf("\nTime= %f",end-start);
    return 0;
}
```

Parallel Code :

```
#include "mpi.h"

#include <stdio.h>

#include <stdlib.h>


#define ROWS_A 500                /* number of rows in matrix A */

#define COLUMNS_A 300            /* number of columns in matrix A */

#define COLUMNS_B 150           /* number of columns in matrix B */

#define MASTER 0                  /* taskid of first task */

#define FROM_MASTER 1             /* setting a message type */

#define FROM_WORKER 2            /* setting a message type */


int main (int argc, char *argv[])

{

int    numtasks,                  /* number of tasks in partition */

    taskid,                      /* a task identifier */

    numworkers,                  /* number of worker tasks */

    source,                      /* task id of message source */

    dest,                        /* task id of message destination */

    mtype,                       /* message type */

    rows,                        /* rows of matrix A sent to each worker */
```

```
averow, extra, offset, /* used to determine rows sent to each worker */

i, j, k, rc;          /* misc */

long double a[ROWS_A][COLUMNS_A],          /* matrix A to be
multiplied */

b[COLUMNS_A][COLUMNS_B],          /* matrix B to be multiplied */

c[ROWS_A][COLUMNS_B];          /* result matrix C */

MPI_Status status;

double start,end;


MPI_Init(&argc,&argv);
start=MPI_Wtime();

MPI_Comm_rank(MPI_COMM_WORLD,&taskid);

MPI_Comm_size(MPI_COMM_WORLD,&numtasks);

if (numtasks < 2 ) {

    printf("Need at least two MPI tasks. Quitting...\n");

    MPI_Abort(MPI_COMM_WORLD, rc);

    exit(1);

}

numworkers = numtasks-1;
```

```
//----- master task
-----//

if (taskid == MASTER)

{

    //printf("mpi_mm has started with %d tasks.\n", numtasks);

    //printf("Initializing arrays...\n");

    for (i=0; i<ROWS_A; i++)

        for (j=0; j<COLUMNS_A; j++)

            a[i][j] = (i+j)*2.658;

    for (i=0; i<COLUMNS_A; i++)

        for (j=0; j<COLUMNS_B; j++)

            b[i][j] = (i+j)*6.658;

    /* Send matrix data to the worker tasks */

    averow = ROWS_A/numworkers;

    extra = ROWS_A%numworkers;

    offset = 0;

    mtype = FROM_MASTER;
```

```
for (dest=1; dest<=numworkers; dest++)

{

    rows = (dest <= extra) ? averow+1 : averow;

    //printf("Sending %d rows to task %d\n",rows,dest,offset);

    MPI_Send(&offset, 1, MPI_INT, dest, mtype, MPI_COMM_WORLD);

    MPI_Send(&rows, 1, MPI_INT, dest, mtype, MPI_COMM_WORLD);

    MPI_Send(&a[offset][0], rows*COLUMNS_A, MPI_LONG_DOUBLE, dest,
mtype,

        MPI_COMM_WORLD);

    MPI_Send(&b, COLUMNS_A*COLUMNS_B, MPI_LONG_DOUBLE, dest, mtype,
MPI_COMM_WORLD);

    offset = offset + rows;

}

/* Receive results from worker tasks */

mtype = FROM_WORKER;

for (i=1; i<=numworkers; i++)

{

    source = i;

    MPI_Recv(&offset, 1, MPI_INT, source, mtype, MPI_COMM_WORLD,
&status);
```

```
        MPI_Recv(&rows, 1, MPI_INT, source, mtype, MPI_COMM_WORLD,
&status);

        MPI_Recv(&c[offset][0], rows*COLUMNS_B, MPI_LONG_DOUBLE, source,
mtype,

                MPI_COMM_WORLD, &status);

        printf("Received results from task %d\n",source);

    }

    /* Print results */

    printf("-----\n");

    printf("Result Matrix:\n");

    for (i=0; i<ROWS_A; i++)

    {

        printf("\n");

        for (j=0; j<COLUMNS_B; j++)

            printf("%6.2Lf    ", c[i][j]);

    }

    printf("\n-----\n");

    printf ("Done.\n");

    end=MPI_Wtime();

    printf("\nTime= %f",end-start);
```



```
}

//----- worker task
-----//

if (taskid > MASTER)

{

    mtype = FROM_MASTER;

    MPI_Recv(&offset, 1, MPI_INT, MASTER, mtype, MPI_COMM_WORLD,
&status);

    MPI_Recv(&rows, 1, MPI_INT, MASTER, mtype, MPI_COMM_WORLD, &status);

    MPI_Recv(&a, rows*COLUMNS_A, MPI_LONG_DOUBLE, MASTER, mtype,
MPI_COMM_WORLD, &status);

    MPI_Recv(&b, COLUMNS_A*COLUMNS_B, MPI_LONG_DOUBLE, MASTER, mtype,
MPI_COMM_WORLD, &status);

    for (k=0; k<COLUMNS_B; k++)

        for (i=0; i<rows; i++)

        {

            c[i][k] = 0.0;

            for (j=0; j<COLUMNS_A; j++)
```

```
        c[i][k] = c[i][k] + a[i][j] * b[j][k];

    }

    mtype = FROM_WORKER;

    MPI_Send(&offset, 1, MPI_INT, MASTER, mtype, MPI_COMM_WORLD);

    MPI_Send(&rows, 1, MPI_INT, MASTER, mtype, MPI_COMM_WORLD);

    MPI_Send(&c, rows*COLUMNS_B, MPI_LONG_DOUBLE, MASTER, mtype,
MPI_COMM_WORLD);

}


MPI_Finalize();

}
```


Observations :

Processes(n)	Runtime	Speedup (s)
1	0.727793	1
2	1.173708	0.6200801221
4	1.625658	0.4476913348
6	1.558417	0.4670078676
8	1.795315	0.4053845704
10	1.772297	0.4106495695
12	1.78564	0.4075810354
16	3.73724	0.1947407713
32	8.186064	0.08890634131
64	23.777508	0.0306084641
128	91.022718	0.007995729154

Runtime vs. Processes(n)



Speedup (s) vs. Processes(n)



Inference: (**Note:** Execution time, graph, and inference will be based on hardware configuration)

- Since MPI is a distributed memory architecture, the communication overhead between nodes causes the parallel code to run slower compared to serial code (running in 1 node or only in master).