

HPC LAB - Vector addition MPI

Name: Paleti Krishnasai

Roll No: CED18I039

Programming Environment: MPI

Problem: Vector addition

Date: 21-10-2021

Hardware Configuration:

PU NAME: Intel(R) Core(TM) i5-8300H CPU @ 2.30GHz

Number of Sockets: 1

Cores per Socket: 4

Threads per core: 2

L1d cache: 128 KiB

L1i cache: 128 KiB

L2 cache: 1 MiB

L3 cache: 8 MiB

```
paleti@paleti-Lenovo-Ideapad-330-15ICH:~$ lscpu
CPU name:      Intel(R) Core(TM) i5-8300H CPU @ 2.30GHz
CPU type:      Intel Cofreelake processor
CPU stepping:  10
*****
Hardware Thread Topology
*****
Sockets:      1
Cores per socket: 4
Threads per core: 2
*****
HWThread      Thread      Core      Socket      Available
0              0          0          0          *
1              0          1          0          *
2              0          2          0          *
3              0          3          0          *
4              1          0          0          *
5              1          1          0          *
6              1          2          0          *
7              1          3          0          *
*****
Socket 0:      ( 0 4 1 5 2 6 3 7 )
*****
Cache Topology
*****
Level:         1
Size:          32 kB
Cache groups:  ( 0 4 ) ( 1 5 ) ( 2 6 ) ( 3 7 )
*****
Level:         2
Size:          256 kB
Cache groups:  ( 0 4 ) ( 1 5 ) ( 2 6 ) ( 3 7 )
*****
Level:         3
Size:          8 MB
Cache groups:  ( 0 4 1 5 2 6 3 7 )
*****
NUMA Topology
*****
NUMA domains:  1
*****
Domain:        0
Processors:    ( 0 1 2 3 4 5 6 7 )
Distances:     10
Free memory:   3546.2 MB
Total memory:  7831.84 MB
*****
```

```
*****
Graphical Topology
*****
Socket 0:
+-----+-----+-----+-----+
| 0 4 | | 1 5 | | 2 6 | | 3 7 |
+-----+-----+-----+-----+
| 32 kB | | 32 kB | | 32 kB | | 32 kB |
+-----+-----+-----+-----+
| 256 kB | | 256 kB | | 256 kB | | 256 kB |
+-----+-----+-----+-----+
|                                     |
|                                     | 8 MB |
+-----+-----+-----+-----+
```

No of nodes: 12 (4 for each as written in machine file).

Serial Code:

```
#include <mpi.h>
#include <stdio.h>
#include <stdlib.h>

#define VSIZE 1000
#define MASTER 0
#define FROM_MASTER 1
#define FROM_WORKER 2

//serial code
int main(int argc, char *argv[]){
    double start,end;
    MPI_Init(&argc, &argv);
    start=MPI_Wtime();
    int i;
    long double a[VSIZE], b[VSIZE], c[VSIZE];
    for (i = 0; i < VSIZE; i++){
        a[i] = i * 156.678;
    }
    for (i = 0; i < VSIZE; i++){
        b[i] = i * 2.0078;
    }
    for (i = 0; i < VSIZE; i++)
        c[i] = a[i] + b[i];
    printf("\nResultant Vector:\n");

    for (i = 0; i < VSIZE; i++)
        printf("\n%Lf +%Lf = %Lf    ", a[i],b[i],c[i]);

    printf("\nFinished.\n");
    end=MPI_Wtime();
    printf("\nTime= %f",end-start);
    return 0;
}
```

Parallel Code: [Point to Point]

```
#include <mpi.h>

#include <stdio.h>

#include <stdlib.h>

#define VSIZE 1000

#define MASTER 0

#define FROM_MASTER 1

#define FROM_WORKER 2

//cluster /parallel code

int main(int argc, char *argv[])

{

    int numtasks, taskid, numworkers, source, dest, mtype, segment,
aveseg, extra, offset, i, j, k, rc;

    long double a[VSIZE], b[VSIZE], c[VSIZE];

    MPI_Status status;

    double start,end;

    MPI_Init(&argc, &argv);

    start=MPI_Wtime();

    MPI_Comm_rank(MPI_COMM_WORLD, &taskid);
```

```
MPI_Comm_size(MPI_COMM_WORLD, &numtasks);

if (numtasks < 2)

{

    printf("Need atleast two MPI tasks. Quitting...\n");

    MPI_Abort(MPI_COMM_WORLD, rc);

    exit(1);

}


char pro_name[MPI_MAX_PROCESSOR_NAME];

int name_len;

MPI_Get_processor_name(pro_name, &name_len);


printf("-From from  %s, rank %d, out of %d\n", pro_name, taskid, numtasks);

numworkers = numtasks - 1;


//master task:


if (taskid == MASTER)

{
```

```
for (i = 0; i < VSIZE; i++) {  
  
    a[i] = i * 156.678;  
  
}  
  
for (i = 0; i < VSIZE; i++) {  
  
    b[i] = i * 2.0078;  
  
}  
  
  
aveseg = VSIZE / numworkers;  
  
extra = VSIZE % numworkers;  
  
offset = 0;  
  
mtype = FROM_MASTER;  
  
for (dest = 1; dest <= numworkers; dest++)  
  
{  
  
    segment = (dest <= extra) ? aveseg + 1 : aveseg;  
  
  
  
    MPI_Send(&offset, 1, MPI_INT, dest, mtype, MPI_COMM_WORLD);  
  
    MPI_Send(&segment, 1, MPI_INT, dest, mtype, MPI_COMM_WORLD);  
  
    MPI_Send(&a[offset], segment, MPI_LONG_DOUBLE, dest, mtype,  
MPI_COMM_WORLD);  
  
    MPI_Send(&b[offset], segment, MPI_LONG_DOUBLE, dest, mtype,  
MPI_COMM_WORLD);  
  
}
```

```
        offset = offset + segment;

    }

    //receive from worker:

    mtype = FROM_WORKER;

    for (i = 1; i <= numworkers; i++)

    {

        source = i;

        MPI_Recv(&offset, 1, MPI_INT, source, mtype, MPI_COMM_WORLD,
&status);

        MPI_Recv(&segment, 1, MPI_INT, source, mtype, MPI_COMM_WORLD,
&status);

        MPI_Recv(&c[offset], segment, MPI_LONG_DOUBLE, source, mtype,
MPI_COMM_WORLD, &status);

    }

    printf("\nResultant Vector:\n");

    for (i = 0; i < VSIZE; i++)

        printf("\n%Lf +%Lf = %Lf    ", a[i],b[i],c[i]);
```

```
    printf("\nFinished.\n");

    end=MPI_Wtime();

    printf("\nTime= %f",end-start);

}

//Worker task:

if (taskid > MASTER)

{

    mtype = FROM_MASTER;

    MPI_Recv(&offset, 1, MPI_INT, MASTER, mtype, MPI_COMM_WORLD,
&status);

    MPI_Recv(&segment, 1, MPI_INT, MASTER, mtype, MPI_COMM_WORLD,
&status);

    MPI_Recv(&a, segment, MPI_LONG_DOUBLE, MASTER, mtype,
MPI_COMM_WORLD, &status);

    MPI_Recv(&b, segment, MPI_LONG_DOUBLE, MASTER, mtype,
MPI_COMM_WORLD, &status);

    //mat addition

    for (i = 0; i < segment; i++)

        c[i] = a[i] + b[i];
```

```
    mtype = FROM_WORKER;

    MPI_Send(&offset, 1, MPI_INT, MASTER, mtype, MPI_COMM_WORLD);

    MPI_Send(&segment, 1, MPI_INT, MASTER, mtype, MPI_COMM_WORLD);

    MPI_Send(&c, segment, MPI_LONG_DOUBLE, MASTER, mtype,
MPI_COMM_WORLD);

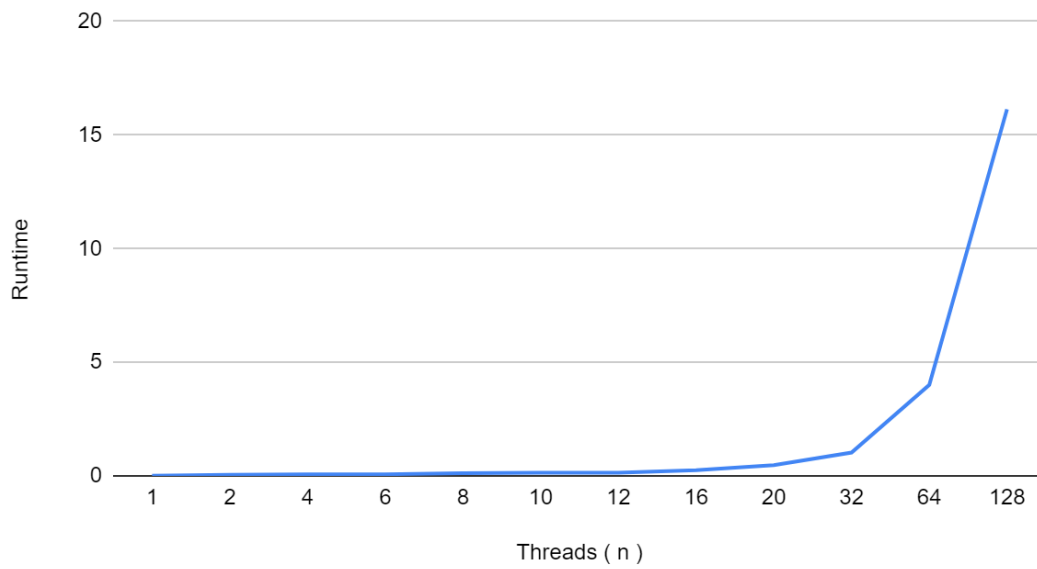
}

MPI_Finalize();
return 0;
}
```


Observations: [Point to Point]

Processes(n)	Runtime	Speedup (s)
1	0.002702	1
2	0.043808	0.06167823229
4	0.068706	0.03932698745
6	0.069162	0.03906769613
8	0.121126	0.02230734937
10	0.136446	0.01980270583
12	0.137429	0.01966106135
16	0.251037	0.01076335361
20	0.470541	0.005742326386
32	1.025241	0.002635477902
64	4.000448	0.0006754243525
128	16.137246	0.0001674387315

Runtime vs. Processes(n)



Speedup (s) vs. Processes(n)



Parallel Code : [Collective]

```
#include <mpi.h>
#include <stdio.h>
#include <math.h>
#include <stdlib.h>

int main(int argc, char *argv[]) {
    int myid, numprocs;
    double startwtime,endwtime;
    int namelen;
    int ARRAY_SIZE = 10000;
    int vector1[ARRAY_SIZE + 50];
    int vector2[ARRAY_SIZE + 50];
    long int vector3[ARRAY_SIZE + 50];
    int i, j;
    int s, s0;
    double totalTime;
    double start,end;
    char processor_name[MPI_MAX_PROCESSOR_NAME];
    MPI_Init(&argc,&argv);
    start=MPI_Wtime();
    MPI_Comm_size(MPI_COMM_WORLD,&numprocs);
    MPI_Comm_rank(MPI_COMM_WORLD,&myid);
    MPI_Get_processor_name(processor_name,&namelen);
    fprintf(stderr,"Process %d on %s\n", myid, processor_name);
    fflush(stderr);
    // Vector 1 Reading
    for (i=0; i < ARRAY_SIZE; i++)
        vector1[i] =i * 156.678;
    // Vector 2 Reading
    for (i=0; i < ARRAY_SIZE; i++)
        vector2[i] =i * 2.0078;

    if(myid == 0)
    {
        MPI_Bcast(&ARRAY_SIZE, 1, MPI_INT, 0, MPI_COMM_WORLD);
        s = (int) floor(ARRAY_SIZE/numprocs);
        s0 = ARRAY_SIZE%numprocs;
        int vector1Receive[s];
        int vector2Receive[s];
        long int vector3Receive[s];
```

```
    if (s0 != 0)
    {
        s = s + 1;
        for(i=0; i < ((s * numprocs) - ARRAY_SIZE); i++)
        {
            vector1[ARRAY_SIZE + i] = 1;
            vector2[ARRAY_SIZE + i] = 1;
        }
    }

    MPI_Bcast(&s, 1, MPI_INT, 0, MPI_COMM_WORLD);
    MPI_Scatter(vector1, s, MPI_INT, vector1Receive, s, MPI_INT, 0,
MPI_COMM_WORLD);
    MPI_Scatter(vector2, s, MPI_INT, vector2Receive, s, MPI_INT, 0,
MPI_COMM_WORLD);
    for(i=0; i<s; i++)
    {
        vector3Receive[i] = vector1Receive[i] + vector2Receive[i];
    }
    MPI_Gather(vector3Receive, s, MPI_LONG, vector3, s, MPI_LONG, 0,
MPI_COMM_WORLD);

    for(i=0; i<ARRAY_SIZE; i++)
        printf("%d %d %d %ld\n",i,vector1[i],vector2[i], vector3[i]);
}
else
{
    MPI_Bcast(&ARRAY_SIZE, 1, MPI_INT, 0, MPI_COMM_WORLD);
    MPI_Bcast(&s, 1, MPI_INT, 0, MPI_COMM_WORLD);
    int vector1Receive[s];
    int vector2Receive[s];
    long int vector3Receive[s];
    MPI_Scatter(vector1, s, MPI_INT, vector1Receive, s, MPI_INT, 0,
MPI_COMM_WORLD);
    MPI_Scatter(vector2, s, MPI_INT, vector2Receive, s, MPI_INT, 0,
MPI_COMM_WORLD);
    for(i=0; i<s; i++)
    {
        vector3Receive[i] = vector1Receive[i] + vector2Receive[i];
    }
}
```

```

        MPI_Gather(vector3Receive, s, MPI_LONG, vector3, s, MPI_LONG, 0,
MPI_COMM_WORLD);

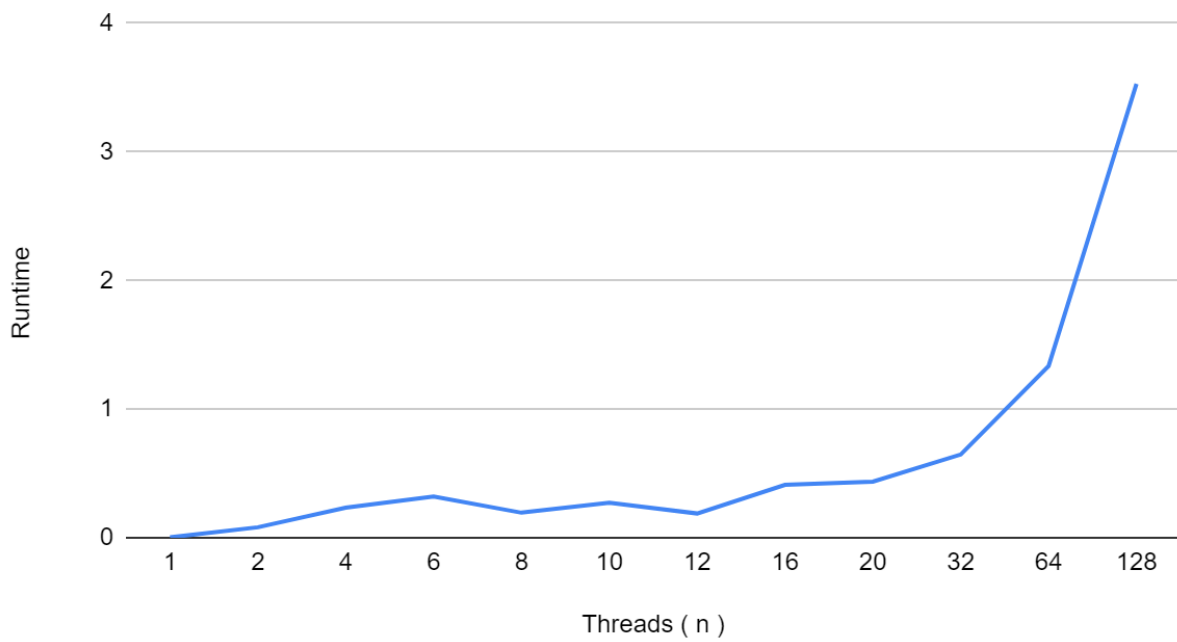
    }
    end=MPI_Wtime();
    printf("\nTime= %f",end-start);
    MPI_Finalize();
}

```


Observations: [Collective]

Processes	Runtime	Speedup (s)
1	0.002702	1
2	0.080808	0.03343728344
4	0.232844	0.01160433595
6	0.320691	0.008425556065
8	0.194362	0.0139018944
10	0.271014	0.009969964651
12	0.188245	0.01435363489
16	0.410867	0.006576337355
20	0.436344	0.00619236199
32	0.646694	0.004178173912
64	1.335289	0.002023531984
128	3.529382	0.000765573123

Runtime vs. Processes(n)



Speedup (s) vs. Processes(n)



Inference: (**Note:** Execution time, graph, and inference will be based on hardware configuration)

- Since MPI is a distributed memory architecture, the communication overhead between nodes causes the parallel code to run slower compared to serial code (running in 1 node or only in master)