

# HPC LAB - Vector Multiplication MPI

**Name:** Paleti Krishnasai

**Roll No:** CED18I039

**Programming Environment:** MPI

**Problem:** Vector Multiplication

**Date:** 21-10-2021

## Hardware Configuration:

PU NAME: Intel(R) Core(TM) i5-8300H CPU @ 2.30GHz

Number of Sockets: 1

Cores per Socket: 4

Threads per core: 2

L1d cache: 128 KiB

L1i cache: 128 KiB

L2 cache: 1 MiB

L3 cache: 8 MiB

```
paleti@paleti-Lenovo-Ideapad-330-15ICH:~$ lllwid-topology
-----
CPU name:      Intel(R) Core(TM) i5-8300H CPU @ 2.30GHz
CPU type:      Intel Coffeelake processor
CPU stepping:  10
-----
Hardware Thread Topology
-----
Sockets:      1
Cores per socket: 4
Threads per core: 2
-----
HWThread      Thread      Core      Socket      Available
-----
0              0              0          0            *
1              0              1          0            *
2              0              2          0            *
3              0              3          0            *
4              1              0          0            *
5              1              1          0            *
6              1              2          0            *
7              1              3          0            *
-----
Socket 0:      ( 0 4 1 5 2 6 3 7 )
-----
Cache Topology
-----
Level:         1
Size:          32 kB
Cache groups:  ( 0 4 ) ( 1 5 ) ( 2 6 ) ( 3 7 )
-----
Level:         2
Size:          256 kB
Cache groups:  ( 0 4 ) ( 1 5 ) ( 2 6 ) ( 3 7 )
-----
Level:         3
Size:          8 MB
Cache groups:  ( 0 4 1 5 2 6 3 7 )
-----
NUMA Topology
-----
NUMA domains:  1
-----
Domain:        0
Processors:    ( 0 1 2 3 4 5 6 7 )
Distances:     10
Free memory:   3546.2 MB
Total memory:  7831.84 MB
-----
```

```
*****
Graphical Topology
*****
Socket 0:
+-----+
| +-----+ +-----+ +-----+ +-----+ |
| | 0 4 | | 1 5 | | 2 6 | | 3 7 | |
| +-----+ +-----+ +-----+ +-----+ |
| | 32 kB | | 32 kB | | 32 kB | | 32 kB | |
| +-----+ +-----+ +-----+ +-----+ |
| | 256 kB | | 256 kB | | 256 kB | | 256 kB | |
| +-----+ +-----+ +-----+ +-----+ |
| | 8 MB | | | | | | | |
| +-----+ +-----+ +-----+ +-----+ |
+-----+
```

No of nodes: 12 ( 4 for each as written in machine file ).

**Serial Code:**

```
#include <mpi.h>
#include <stdio.h>
#include <stdlib.h>

#define VSIZE 1000
#define MASTER 0
#define FROM_MASTER 1
#define FROM_WORKER 2

//serial code
int main(int argc, char *argv[]){
    double start,end;
    MPI_Init(&argc, &argv);
    start=MPI_Wtime();
    int i;
    long double a[VSIZE], b[VSIZE], c[VSIZE];
    for (i = 0; i < VSIZE; i++){
        a[i] = i * 156.678;
    }
    for (i = 0; i < VSIZE; i++){
        b[i] = i * 2.0078;
    }
    for (i = 0; i < VSIZE; i++)
        c[i] = a[i] * b[i];
    printf("\nResultant Vector:\n");

    for (i = 0; i < VSIZE; i++)
        printf("\n%Lf +%Lf = %Lf    ", a[i],b[i],c[i]);

    printf("\nFinished.\n");
    end=MPI_Wtime();
    printf("\nTime= %f",end-start);
    return 0;
}
```

### Parallel Code: [ Point to Point ]

```
#include <mpi.h>

#include <stdio.h>

#include <stdlib.h>

#define VSIZE 1000

#define MASTER 0

#define FROM_MASTER 1

#define FROM_WORKER 2

int main(int argc, char *argv[])

{

    int numtasks, taskid, numworkers, source, dest, mtype, segment,
aveseg, extra, offset, i, j, k, rc;

    long double a[VSIZE], b[VSIZE], c[VSIZE];

    MPI_Status status;

    double start,end;

    MPI_Init(&argc, &argv);

    start=MPI_Wtime();

    MPI_Comm_rank(MPI_COMM_WORLD, &taskid);

    MPI_Comm_size(MPI_COMM_WORLD, &numtasks);
```

```
if (numtasks < 2)

{

    printf("Need atleast two MPI tasks. Quitting...\n");

    MPI_Abort(MPI_COMM_WORLD, rc);

    exit(1);

}


char pro_name[MPI_MAX_PROCESSOR_NAME];

int name_len;

MPI_Get_processor_name(pro_name, &name_len);


printf("-From from  %s, rank %d, out of %d\n", pro_name, taskid, numtasks);

numworkers = numtasks - 1;


//master task:


if (taskid == MASTER)

{

    for (i = 0; i < VSIZE; i++){
```

```
        a[i] = i * 156.678;

    }

    for (i = 0; i < VSIZE; i++){

        b[i] = i * 2.0078;

    }


    aveseq = VSIZE / numworkers;

    extra = VSIZE % numworkers;

    offset = 0;

    mtype = FROM_MASTER;

    for (dest = 1; dest <= numworkers; dest++)

    {

        segment = (dest <= extra) ? aveseq + 1 : aveseq;

        MPI_Send(&offset, 1, MPI_INT, dest, mtype, MPI_COMM_WORLD);

        MPI_Send(&segment, 1, MPI_INT, dest, mtype, MPI_COMM_WORLD);

        MPI_Send(&a[offset], segment, MPI_LONG_DOUBLE, dest, mtype,
MPI_COMM_WORLD);

        MPI_Send(&b[offset], segment, MPI_LONG_DOUBLE, dest, mtype,
MPI_COMM_WORLD);

        offset = offset + segment;

    }

}
```

```
}

//receive from worker:

mtype = FROM_WORKER;

for (i = 1; i <= numworkers; i++)

{

    source = i;

    MPI_Recv(&offset, 1, MPI_INT, source, mtype, MPI_COMM_WORLD,
&status);

    MPI_Recv(&segment, 1, MPI_INT, source, mtype, MPI_COMM_WORLD,
&status);

    MPI_Recv(&c[offset], segment, MPI_LONG_DOUBLE, source, mtype,
MPI_COMM_WORLD, &status);

}

printf("\nResultant Vector:\n");

for (i = 0; i < VSIZE; i++)

    printf("\n%Lf +%Lf = %Lf    ", a[i],b[i],c[i]);

printf("\nFinished.\n");
```

```
end=MPI_Wtime();

printf("\nTime= %f",end-start);

}

//Worker task:

if (taskid > MASTER)

{

    mtype = FROM_MASTER;

    MPI_Recv(&offset, 1, MPI_INT, MASTER, mtype, MPI_COMM_WORLD,
&status);

    MPI_Recv(&segment, 1, MPI_INT, MASTER, mtype, MPI_COMM_WORLD,
&status);

    MPI_Recv(&a, segment, MPI_LONG_DOUBLE, MASTER, mtype,
MPI_COMM_WORLD, &status);

    MPI_Recv(&b, segment, MPI_LONG_DOUBLE, MASTER, mtype,
MPI_COMM_WORLD, &status);

    //mat multiplication

    for (i = 0; i < segment; i++)

        c[i] = a[i] * b[i];
```

```
    mtype = FROM_WORKER;

    MPI_Send(&offset, 1, MPI_INT, MASTER, mtype, MPI_COMM_WORLD);

    MPI_Send(&segment, 1, MPI_INT, MASTER, mtype, MPI_COMM_WORLD);

    MPI_Send(&c, segment, MPI_LONG_DOUBLE, MASTER, mtype,
MPI_COMM_WORLD);

}


MPI_Finalize();
return 0;
}
```

#### Observations: [ Point to Point ]

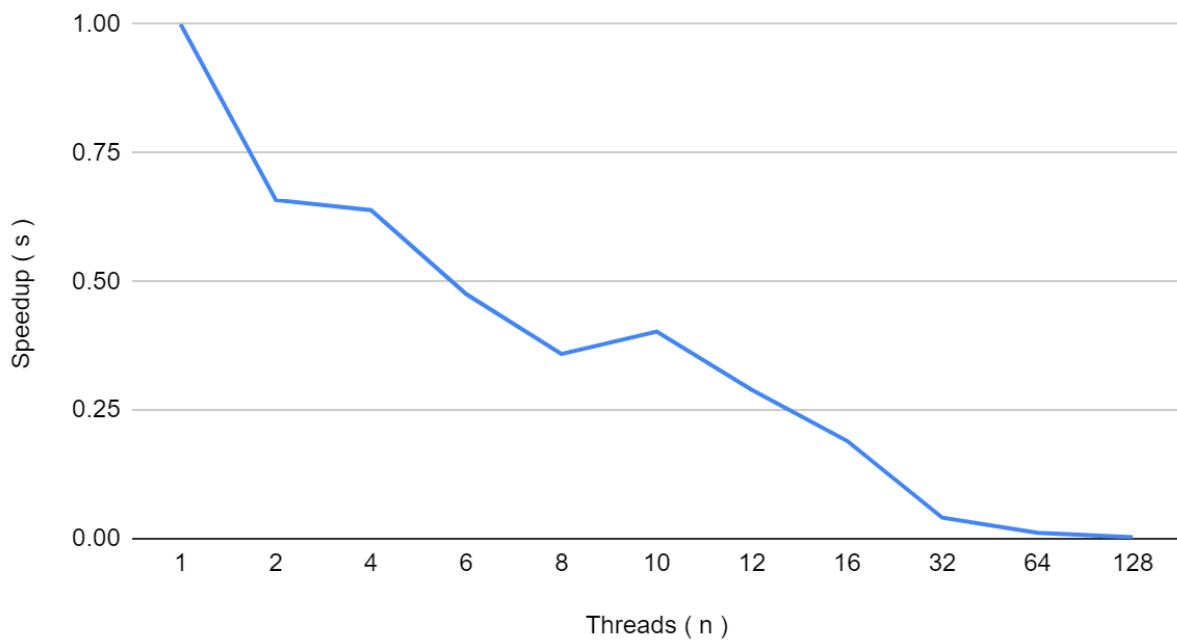
Processes( n )	Runtime	Speedup ( s )
1	0.042741	1
2	0.064913	0.6584351363
4	0.066917	0.6387166191
6	0.089911	0.4753700882
8	0.119081	0.3589237578
10	0.106206	0.4024348907
12	0.147787	0.2892067638
16	0.2249	0.1900444642
32	1.046669	0.04083525928
64	3.7924277	0.01127008961
128	15.069948	0.002836174352



## Runtime vs. Processes



## Speedup ( s ) vs. Processes( n )



**Parallel Code: [ Collective ]**

```
#include "mpi.h"
#include <stdio.h>
#include <math.h>
#include <stdlib.h>

int main(int argc, char *argv[]) {
    int myid, numprocs;
    double startwtime,
    endwtime;
    int namelen;
    int ARRAY_SIZE = 100000;
    int vector1[ARRAY_SIZE];
    int vector2[ARRAY_SIZE];
    int vector3[ARRAY_SIZE];
    int i, j;
    int s, s0;
    double totalTime;
    double start, end;
    char processor_name[MPI_MAX_PROCESSOR_NAME];
    MPI_Init(&argc, &argv);
    start = MPI_Wtime();
    MPI_Comm_size(MPI_COMM_WORLD, &numprocs);
    MPI_Comm_rank(MPI_COMM_WORLD, &myid);
    MPI_Get_processor_name(processor_name, &namelen);
    fprintf(stderr, "Process %d on %s\n", myid, processor_name);
    fflush(stderr);
    // Vector 1 Reading
    for (i=0; i < ARRAY_SIZE; i++)
        vector1[i] = i * 156.678;
    // Vector 2 Reading
    for (i=0; i < ARRAY_SIZE; i++)
        vector2[i] = i * 2.0078;

    if(myid == 0)
    {
        MPI_Bcast(&ARRAY_SIZE, 1, MPI_INT, 0, MPI_COMM_WORLD);
        s = (int) floor(ARRAY_SIZE/numprocs);
        s0 = ARRAY_SIZE%numprocs;
        int vector1Receive[s];
        int vector2Receive[s];
```

```
int vector3Receive[s];
if (s0 != 0)
{
    s = s + 1;
    for(i=0; i < ((s * numprocs) - ARRAY_SIZE); i++)
    {
        vector1[ARRAY_SIZE + i] = 0;
        vector2[ARRAY_SIZE + i] = 0;
    }
}
MPI_Bcast(&s, 1, MPI_INT, 0, MPI_COMM_WORLD);
MPI_Scatter(vector1, s, MPI_INT, vector1Receive, s, MPI_INT, 0,
MPI_COMM_WORLD);
MPI_Scatter(vector2, s, MPI_INT, vector2Receive, s, MPI_INT, 0,
MPI_COMM_WORLD);
for(i=0; i<s; i++)
{
    vector3Receive[i] = vector1Receive[i] * vector2Receive[i];
}
MPI_Gather(vector3Receive, s, MPI_INT, vector3, s, MPI_INT, 0,
MPI_COMM_WORLD);


for(i=0; i<ARRAY_SIZE; i++)
    printf("%d\n", vector3[i]);
}
else
{
    MPI_Bcast(&ARRAY_SIZE, 1, MPI_INT, 0, MPI_COMM_WORLD);
    MPI_Bcast(&s, 1, MPI_INT, 0, MPI_COMM_WORLD);
    int vector1Receive[s];
    int vector2Receive[s];
    int vector3Receive[s];
    MPI_Scatter(vector1, s, MPI_INT, vector1Receive, s, MPI_INT, 0,
MPI_COMM_WORLD);
    MPI_Scatter(vector2, s, MPI_INT, vector2Receive, s, MPI_INT, 0,
MPI_COMM_WORLD);
    for(i=0; i<s; i++)
    {
        vector3Receive[i] = vector1Receive[i] * vector2Receive[i];
    }
}
```

```
MPI_Gather(vector3Receive, s, MPI_INT, vector3, s, MPI_INT, 0,  
MPI_COMM_WORLD);  
  
}  
end=MPI_Wtime();  
printf("\nTime= %f",end-start);  
MPI_Finalize();  
}
```


**Observations [ Collective ] :**

Processes( n )	Runtime	Speedup ( s )
1	0.042741	1
2	0.73142	0.05843564573
4	1.977666	0.02161183941
6	2.065529	0.02069251993
8	2.099447	0.02035821814
10	2.279888	0.01874697354
12	2.240546	0.01907615376
16	3.053671	0.01399659623
32	4.816833	0.008873257595
64	8.681692	0.004923118673
128	16.589044	0.002576459499

## Runtime vs. Processes



## Speedup ( s ) vs. Processes( n )



**Inference:** (**Note:** Execution time, graph, and inference will be based on hardware configuration)

- Since MPI is a distributed memory architecture, the communication overhead between nodes causes the parallel code to run slower compared to serial code ( running in 1 node or only in master )