

# HPC LAB - Vector Addition CUDA

**Name:** Paleti Krishnasai

**Roll No: CED18I039**

### Programming Environment: CUDA (collab)

### Problem: Vector Addition

### Hardware Configuration:

PU NAME: Intel(R) Core(TM) i5-8300H CPU @ 2.30GHz

Number of Sockets: 1

Cores per Socket: 4

Threads per core: 2

L1d cache: 128 KiB

L1i cache: 128 KiB

L2 cache: 1 MiB

L3 cache: 8 MiB

```

palett@palett-Lenovo-ideapad-330-15ICH: ~$ lscpu
CPU name:      Intel(R) Core(TM) i5-8300H CPU @ 2.30GHz
CPU type:      Intel Coffeelake processor
CPU stepping:  10
*****
Hardware Thread Topology
*****
Sockets:       1
Cores per socket: 4
Threads per core: 2
*****
HWThread      Thread      Core      Socket      Available
-----
0              0              0          0            *
1              0              1          0            *
2              0              2          0            *
3              0              3          0            *
4              1              0          0            *
5              1              1          0            *
6              1              2          0            *
7              1              3          0            *
*****
Socket 0:      ( 0 4 1 5 2 6 3 7 )
*****
Cache Topology
*****
Level:         1
Size:          32 kB
Cache groups:  ( 0 4 ) ( 1 5 ) ( 2 6 ) ( 3 7 )
*****
Level:         2
Size:          256 kB
Cache groups:  ( 0 4 ) ( 1 5 ) ( 2 6 ) ( 3 7 )
*****
Level:         3
Size:          8 MB
Cache groups:  ( 0 4 1 5 2 6 3 7 )
*****
NUMA Topology
*****
NUMA domains:  1
*****
Domain:        0
Processors:    ( 0 1 2 3 4 5 6 7 )
Distances:     10
Free memory:   3546.2 MB
Total memory:  7831.84 MB

```

```

*****
Graphical Topology
*****
Socket 0:
+-----+
| +-----+ +-----+ +-----+ +-----+ |
| | 0 4 | | 1 5 | | 2 6 | | 3 7 | |
| +-----+ +-----+ +-----+ +-----+ |
| +-----+ +-----+ +-----+ +-----+ |
| | 32 kB | | 32 kB | | 32 kB | | 32 kB | |
| +-----+ +-----+ +-----+ +-----+ |
| +-----+ +-----+ +-----+ +-----+ |
| | 256 kB | | 256 kB | | 256 kB | | 256 kB | |
| +-----+ +-----+ +-----+ +-----+ |
| | 8 MB | | | | | | |
| +-----+ +-----+ +-----+ +-----+ |
+-----+

```

### CUDA code:

```
%%cu
#include <stdio.h>
#include <stdlib.h>
#include<stdlib.h>
#include<time.h>
#include<math.h>
#define n_size 100000
__global__ void add(float *a, float *b, float *c) {
int index=threadIdx.x+blockIdx.x*blockDim.x;
c[index]=a[index]+b[index];
}
void random_init(float a[],int ch)
{
    srand(time(NULL));
    if(ch==0)
    {
        for(int i=0;i<n_size;i++)
        {
            a[i]=((float)rand()/(float)(RAND_MAX)) * 5.0;
        }
    }
    else
    {
        for(int i=0;i<n_size;i++)
        {
            a[i]=(i+1);
        }
    }
}
int main() {
    float a[n_size], b[n_size],c[n_size];
    cudaEvent_t start, end;
    // host copies of variables a, b & c
    float *d_a, *d_b, *d_c;
    // device copies of variables a, b & c
    int size = n_size*sizeof(float);
    // Allocate space for device copies of a, b, c
    cudaMalloc((void **)&d_a, size);
```

```
cudaMalloc((void **)&d_b, size);
cudaMalloc((void **)&d_c, size);
// Create Event for time
cudaEventCreate(&start);
cudaEventCreate(&end);
// Setup input values
random_init(a,0);
random_init(b,0);
// Copy inputs to device
cudaMemcpy(d_a, &a, size, cudaMemcpyHostToDevice);
cudaMemcpy(d_b, &b, size, cudaMemcpyHostToDevice);
int Thread[]={1,2,4,6,8,10,12,16,20,32,64,128,150};
int thread_arr_size=13;
for(int i=0;i<thread_arr_size;i++)
{
    int Threads=Thread[i];
    cudaEventRecord(start);
    // Launch add() kernel on GPU
    add<<<n_size/Threads,Threads>>>(d_a, d_b, d_c);
    cudaEventRecord(end);
    cudaEventSynchronize(end);
    float time = 0;
    cudaEventElapsedTime(&time, start, end);
    // Copy result back to host
    cudaError err = cudaMemcpy(&c, d_c, size, cudaMemcpyDeviceToHost);
    if(err!=cudaSuccess) {
        printf("CUDA error copying to Host: %s\n",
        cudaGetErrorString(err));
    }
    int flag=0;
    for(int i=0;i<n_size;i++)
    {
        if(c[i]!=(a[i]+b[i]))
        {
            flag=1;
            break;
        }
    }
    if(flag==0)
    {
```

```
        printf("Program Executed as Expected\n");
        printf("Time Taken by the program for %d
Threads=%f\n",Threads,time);
    }
    else
    {
        printf("Vector Addition hasnt been done properly,Mismatch in
Values!!!\n");
    }

}
// Cleanup
cudaFree(d_a);
cudaFree(d_b);
cudaFree(d_c);
return 0;
}
```

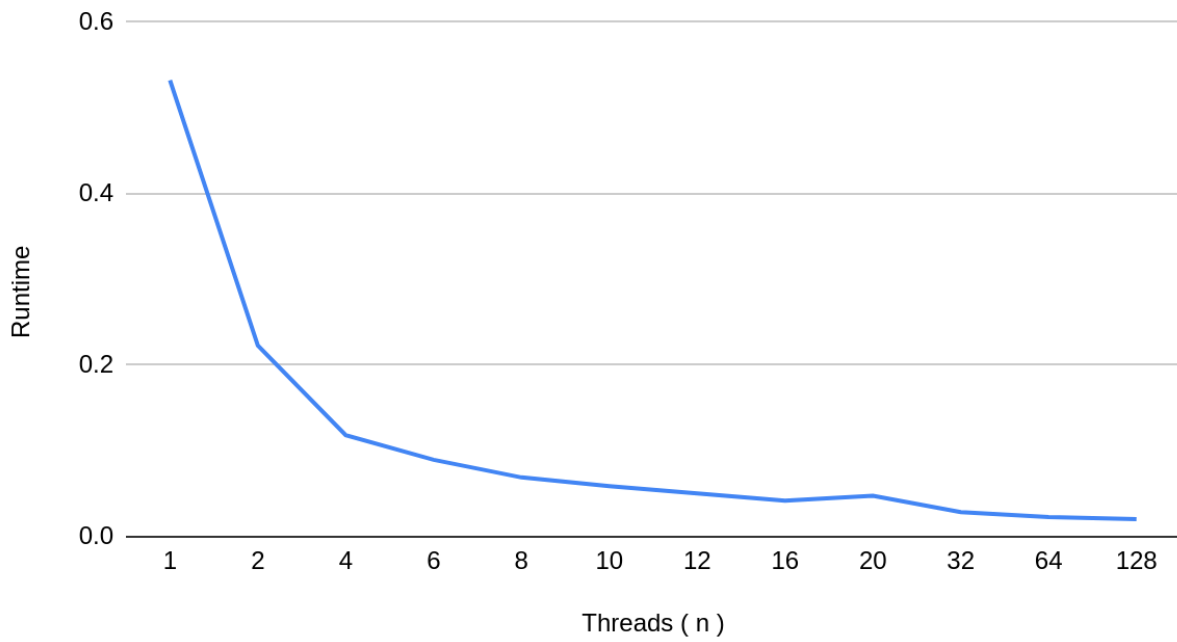
**Observations:**

The function add() is executed for an array of size 'N' on GPU for "N/T" times for 'T' Threads.

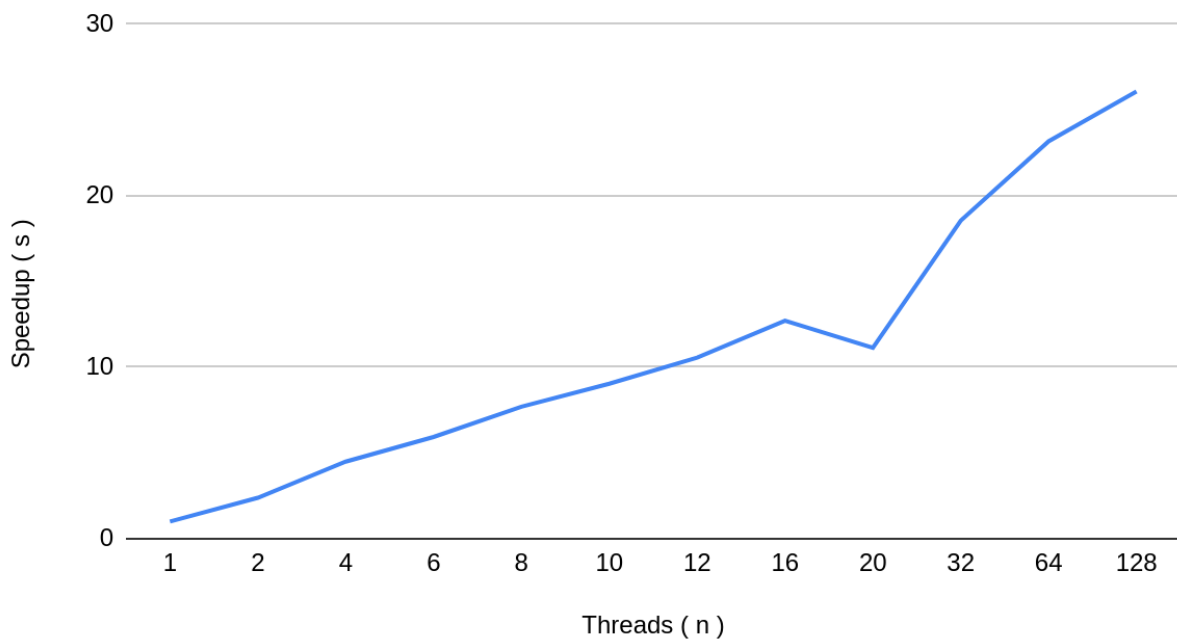
Blocks	Threads ( n )	Runtime	Speedup ( s )	Parallelization Fraction
N	1	0.532512	1	
N/2	2	0.22304	2.387517934	1.162309957
N/4	4	0.118464	4.495137763	1.036716543
N/6	6	0.089728	5.934736091	0.9978006129
N/8	8	0.069152	7.700601573	0.9944457321
N/10	10	0.058912	9.03910918	0.9881884769
N/12	12	0.050432	10.55901015	0.9875936214
N/16	16	0.041888	12.71275783	0.9827614526
N/20	20	0.04784	11.13110368	0.9580648936
N/32	32	0.028704	18.55183946	0.9766162471
N/64	64	0.022976	23.17688022	0.9720417061
N/128	128	0.020416	26.0830721	0.9692330914

### Graphical Inference:

Runtime vs. Threads ( n )



Speedup ( s ) vs. Threads ( n )



**Inference:**

From the Graph it is observed that the speedup starts to increase steadily after 20 threads constantly and the highest Speedup is for 128 threads. It also increases steadily from 1 to 16 threads, but goes down at 20 and rises steadily again.

Runtime Decreases drastically from 1 to 4 threads and steadily decreases from then on till 128, with a minor rise at threads = 20.