

HPC LAB - Vector Multiplication CUDA

Name: Paleti Krishnasai

Roll No: CED18I039

Programming Environment: CUDA (collab)

Problem: Vector Multiplication

Hardware Configuration:

PU NAME: Intel(R) Core(TM) i5-8300H CPU @ 2.30GHz

Number of Sockets: 1

Cores per Socket: 4

Threads per core: 2

L1d cache: 128 KiB

L1i cache: 128 KiB

L2 cache: 1 MiB

L3 cache: 8 MiB

```
paleti@paleti-Lenovo-Ideapad-330-15ICH:~$ likwid-topology
-----
CPU name:      Intel(R) Core(TM) i5-8300H CPU @ 2.30GHz
CPU type:      Intel Coffeelake processor
CPU stepping:  10
*****
Hardware Thread Topology
*****
Sockets:       1
Cores per socket: 4
Threads per core: 2
-----
HWThread      Thread      Core      Socket      Available
-----
0              0              0          0            *
1              0              1          0            *
2              0              2          0            *
3              0              3          0            *
4              1              0          0            *
5              1              1          0            *
6              1              2          0            *
7              1              3          0            *
-----
Socket 0:      ( 0 4 1 5 2 6 3 7 )
-----
Cache Topology
*****
Level:         1
Size:          32 kB
Cache groups:  ( 0 4 ) ( 1 5 ) ( 2 6 ) ( 3 7 )
-----
Level:         2
Size:          256 kB
Cache groups:  ( 0 4 ) ( 1 5 ) ( 2 6 ) ( 3 7 )
-----
Level:         3
Size:          8 MB
Cache groups:  ( 0 4 1 5 2 6 3 7 )
-----
*****
NUMA Topology
*****
NUMA domains:  1
-----
Domain:        0
Processors:    ( 0 1 2 3 4 5 6 7 )
Distances:     10
Free memory:   3546.2 MB
Total memory:  7831.84 MB
-----
```

```
*****
Graphical Topology
*****
Socket 0:
+-----+
| +-----+ +-----+ +-----+ +-----+ |
| | 0 4 | | 1 5 | | 2 6 | | 3 7 | |
| +-----+ +-----+ +-----+ +-----+ |
| | 32 kB | | 32 kB | | 32 kB | | 32 kB | |
| +-----+ +-----+ +-----+ +-----+ |
| | 256 kB | | 256 kB | | 256 kB | | 256 kB | |
| +-----+ +-----+ +-----+ +-----+ |
| |                                     | |
| |                               8 MB | |
| +-----+ +-----+ +-----+ +-----+ |
+-----+
```

CUDA code:

```
//VEctor Mul
%%cu
#include <stdio.h>
#include <stdlib.h>
#include<stdlib.h>
#include<time.h>
#include<math.h>
#include<unistd.h>
//#define Threads 2
#define n_size 100000
__global__ void mul(float *a, float *b, float *c) {
int index=threadIdx.x+blockIdx.x*blockDim.x;
c[index]=a[index]*b[index];
}
void random_init(float a[],int ch)
{
    srand(time(NULL));
    if(ch==0)
    {
        for(int i=0;i<n_size;i++)
        {
            a[i]=((float)rand()/(float)(RAND_MAX)) * 5.0;
        }
    }
    else
    {
        for(int i=0;i<n_size;i++)
        {
            a[i]=(i+1);
        }
    }
}
int main() {
    float a[n_size], b[n_size],c[n_size];
    cudaEvent_t start, end;
    // host copies of variables a, b & c
    float *d_a, *d_b, *d_c;
    // device copies of variables a, b & c
```

```
int size = n_size*sizeof(float);
// Allocate space for device copies of a, b, c
cudaMalloc((void **)&d_a, size);
cudaMalloc((void **)&d_b, size);
cudaMalloc((void **)&d_c, size);
// Create Event for time
cudaEventCreate(&start);
cudaEventCreate(&end);
// Setup input values
random_init(a,0);
random_init(b,0);
// Copy inputs to device
cudaMemcpy(d_a, &a, size, cudaMemcpyHostToDevice);
cudaMemcpy(d_b, &b, size, cudaMemcpyHostToDevice);
int Thread[]={1,2,4,6,8,10,12,16,20,32,64,128,150};
int thread_arr_size=13;
for(int i=0;i<thread_arr_size;i++)
{
    sleep(1);
    int Threads=Thread[i];
    cudaEventRecord(start);
    // Launch add() kernel on GPU
    mul<<<n_size/Threads,Threads>>>(d_a, d_b, d_c);
    cudaEventRecord(end);
    cudaEventSynchronize(end);
    float time = 0;
    cudaEventElapsedTime(&time, start, end);
    // Copy result back to host
    cudaError err = cudaMemcpy(&c, d_c, size, cudaMemcpyDeviceToHost);
    if(err!=cudaSuccess) {
        printf("CUDA error copying to Host: %s\n",
        cudaGetErrorString(err));
    }
    int flag=0;
    for(int i=0;i<n_size;i++)
    {
        //printf("Result[%d]=%f\n",i+1,c[i]);
        if(c[i]!=(a[i]*b[i]))
        {
            flag=1;
            break;
        }
    }
}
```

```
    }  
    if(flag==0)  
    {  
        printf("Program Executed as Expected\n");  
        printf("Time Taken by the program for %d  
Threads=%f\n",Threads,time);  
        //printf("%f\n",time);  
    }  
    else  
    {  
        printf("Vector Addition hasnt been done properly,Mismatch in  
Values!!!\n");  
    }  
  
}  
// Cleanup  
cudaFree(d_a);  
cudaFree(d_b);  
cudaFree(d_c);  
return 0;  
}
```


Observations:

The function add() is executed for an array of size 'N' on GPU for "N/T" times for 'T' Threads.


Blocks	Threads (n)	Runtime	Speedup (s)	Parallelization Fraction
N	1	0.594528	1	
N/2	2	0.236224	2.516797616	1.205339362
N/4	4	0.136448	4.35717636	1.027324757
N/6	6	0.103584	5.739573679	0.9909252382
N/8	8	0.08224	7.229182879	0.9847677485
N/10	10	0.09744	6.101477833	0.9290058668
N/12	12	0.067264	8.838725024	0.967485284
N/16	16	0.059072	10.06446371	0.9606832086
N/20	20	0.066304	8.966698842	0.9352381438
N/32	32	0.05888	10.09728261	0.9300267906
N/64	64	0.037728	15.75826972	0.9514069905
N/128	128	0.035616	16.69272237	0.9474959664

Graphical Inference:

Runtime vs. Threads (n)



Speedup (s) vs. Threads (n)



Inference:

From the Graph, it is observed that the speedup increase in an up and down manner after 8 threads, and the highest Speedup is for 128 threads.

Runtime Decreases drastically from 1 to 4 threads and steadily decreases from then on till 128, with a minor rise at threads = 10.