

HPC LAB - Dot Product MPI

Name: Paleti Krishnasai
Roll No: CED18I039
Programming Environment: MPI
Problem: Dot Product

Hardware Configuration:

PU NAME: Intel(R) Core(TM) i5-8300H CPU @ 2.30GHz
Number of Sockets: 1
Cores per Socket: 4
Threads per core: 2
L1d cache: 128 KiB
L1i cache: 128 KiB
L2 cache: 1 MiB
L3 cache: 8 MiB

```
paleti@paleti-Lenovo-Ideapad-330-15ICH:~$ lscpu
CPU name: Intel(R) Core(TM) i5-8300H CPU @ 2.30GHz
CPU type: Intel Cofeelake processor
CPU stepping: 10
*****
Hardware Thread Topology
*****
Sockets: 1
Cores per socket: 4
Threads per core: 2
*****
HWThread Thread Core Socket Available
0 0 0 0 *
1 0 1 0 *
2 0 2 0 *
3 0 3 0 *
4 1 0 0 *
5 1 1 0 *
6 1 2 0 *
7 1 3 0 *
*****
Socket 0: ( 0 4 1 5 2 6 3 7 )
*****
Cache Topology
*****
Level: 1
Size: 32 kB
Cache groups: ( 0 4 ) ( 1 5 ) ( 2 6 ) ( 3 7 )
*****
Level: 2
Size: 256 kB
Cache groups: ( 0 4 ) ( 1 5 ) ( 2 6 ) ( 3 7 )
*****
Level: 3
Size: 8 MB
Cache groups: ( 0 4 1 5 2 6 3 7 )
*****
NUMA Topology
*****
NUMA domains: 1
*****
Domain: 0
Processors: ( 0 1 2 3 4 5 6 7 )
Distances: 10
Free memory: 3546.2 MB
Total memory: 7831.84 MB
*****
```

```
*****
Graphical Topology
*****
Socket 0:
+-----+
| +-----+ +-----+ +-----+ +-----+ |
| | 0 4 | | 1 5 | | 2 6 | | 3 7 | |
| +-----+ +-----+ +-----+ +-----+ |
| | 32 kB | | 32 kB | | 32 kB | | 32 kB | |
| +-----+ +-----+ +-----+ +-----+ |
| | 256 kB | | 256 kB | | 256 kB | | 256 kB | |
| +-----+ +-----+ +-----+ +-----+ |
| | | | 8 MB | | | |
| +-----+ +-----+ +-----+ +-----+ |
+-----+
```

No of nodes: 12 (4 for each as written in the machine file).

Serial Code:

```
#include "mpi.h"

#include <stdio.h>

#include <stdlib.h>

#include <float.h>

#include <time.h>
#define SIZE 1200000

int main(int argc, char *argv[])

{

    srand(time(0));

    double start,end;

    MPI_Init(&argc,&argv);

    start=MPI_Wtime();

    int i = 0, n = 120000;

    double a[n],b[n], sum = 0;

    for(i=0;i<120000;i++)

    {

        a[i] = (i+1)*4.37 ;

        b[i] = (i+1)*5.36 ;

    }

    for(i=0;i<120000;i++)
```

```
{  
  
    sum += a[i]*b[i];  
  
}  
  
end=MPI_Wtime();  
  
printf("\nTime= %f",end-start);  
  
return 0;  
}
```

Parallel Code:

```
#include "mpi.h"  
#include <stdio.h>  
#include <math.h>  
#include <stdlib.h>  
int main(int argc, char *argv[]) {  
    int myid, numprocs;  
    double start,end;  
    int namelen;  
    int ARRAY_SIZE = 10000;  
    int vector1[ARRAY_SIZE+50];  
    int vector2[ARRAY_SIZE+50];  
    long int vector3[ARRAY_SIZE],part_sum=0,sum=0;  
    int i, j;  
    int s, s0;  
    //double totalTime;  
    char processor_name[MPI_MAX_PROCESSOR_NAME];  
    MPI_Init(&argc,&argv);  
    start=MPI_Wtime();  
    MPI_Comm_size(MPI_COMM_WORLD,&numprocs);  
    MPI_Comm_rank(MPI_COMM_WORLD,&myid);  
    MPI_Get_processor_name(processor_name,&namelen);  
    //fprintf(stderr,"Process %d on %s\n", myid, processor_name);  
    fflush(stderr);
```

```
// Vector 1 Reading
for (i=0; i < ARRAY_SIZE; i++)
    vector1[i] = (i+1)*4.37;
// Vector 2 Reading
for (i=0; i < ARRAY_SIZE; i++)
    vector2[i] = (i+1)*5.36;

if(myid == 0)
{
    MPI_Bcast(&ARRAY_SIZE, 1, MPI_INT, 0, MPI_COMM_WORLD);
    s = (int) floor(ARRAY_SIZE/numprocs);
    s0 = ARRAY_SIZE%numprocs;
    int vector1Receive[s];
    int vector2Receive[s];
    long int vector3Receive[s];
    if (s0 != 0)
    {
        s = s + 1;
        for(i=0; i < ((s * numprocs) - ARRAY_SIZE); i++)
        {
            vector1[ARRAY_SIZE + i] = 0;
            vector2[ARRAY_SIZE + i] = 0;
        }
    }
    MPI_Bcast(&s, 1, MPI_INT, 0, MPI_COMM_WORLD);
    MPI_Scatter(vector1, s, MPI_INT, vector1Receive, s, MPI_INT, 0,
MPI_COMM_WORLD);
    MPI_Scatter(vector2, s, MPI_INT, vector2Receive, s, MPI_INT, 0,
MPI_COMM_WORLD);
    for(i=0; i<s; i++)
    {
        vector3Receive[i] = vector1Receive[i] * vector2Receive[i];
        part_sum += vector3Receive[i];
    }
    MPI_Gather(vector3Receive, s, MPI_LONG, vector3, s, MPI_LONG, 0,
MPI_COMM_WORLD);

    //for(i=0; i<ARRAY_SIZE; i++)
        //printf("%ld\n", vector3[i]);
}
```


```
else
{
    MPI_Bcast(&ARRAY_SIZE, 1, MPI_INT, 0, MPI_COMM_WORLD);
    MPI_Bcast(&s, 1, MPI_INT, 0, MPI_COMM_WORLD);
    int vector1Receive[s];
    int vector2Receive[s];
    long int vector3Receive[s];
    MPI_Scatter(vector1, s, MPI_INT, vector1Receive, s, MPI_INT, 0,
MPI_COMM_WORLD);
    MPI_Scatter(vector2, s, MPI_INT, vector2Receive, s, MPI_INT, 0,
MPI_COMM_WORLD);
    for(i=0; i<s; i++)
    {
        vector3Receive[i] = vector1Receive[i] * vector2Receive[i];
        part_sum += vector3Receive[i];
    }
    MPI_Gather(vector3Receive, s, MPI_LONG, vector3, s, MPI_LONG, 0,
MPI_COMM_WORLD);
}
MPI_Reduce(&part_sum, &sum, 1, MPI_LONG, MPI_SUM, 0, MPI_COMM_WORLD);
if (myid == 0)
{
    printf("\ndot product = %ld \n", sum);

    end=MPI_Wtime();
    printf("\nTime_hello= %f",end-start);
}
MPI_Finalize();
}
```

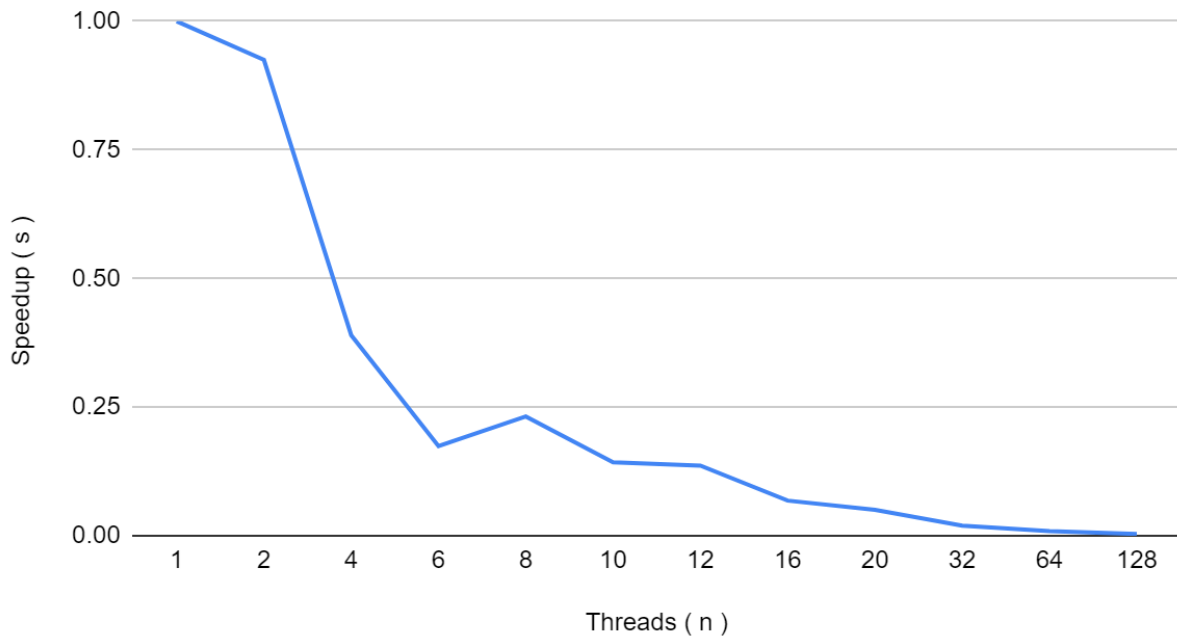
Observations :

Processes (n)	Runtime	Speedup (s)
1	0.010808	1
2	0.011686	0.9248673627
4	0.02773	0.3897583844
6	0.062029	0.1742410808
8	0.046634	0.2317622336
10	0.07572	0.1427363973
12	0.079346	0.1362135457
16	0.158487	0.06819486772
20	0.214793	0.05031821335
32	0.555407	0.0194596035
64	1.212303	0.008915262933
128	3.290684	0.003284423542

Runtime vs. Processes (n)



Speedup (s) vs. Processes (n)



Inference: (Note: Execution time, graph, and inference will be based on hardware configuration)

- Since MPI is a distributed memory architecture, the communication overhead between nodes causes the parallel code to run slower compared to serial code (running in 1 node or only in master).