

HPC LAB - Sum of N Numbers MPI

Name: Paleti Krishnasai

Roll No: CED18I039

Programming Environment: MPI

Problem:Sum of N numbers

Hardware Configuration:

PU NAME: Intel(R) Core(TM) i5-8300H CPU @ 2.30GHz

Number of Sockets: 1

Cores per Socket: 4

Threads per core: 2

L1d cache: 128 KiB

L1i cache: 128 KiB

L2 cache: 1 MiB

L3 cache: 8 MiB

```

# cat /etc/passwd | grep root | wc -l
0
root@kali:~# dmesg | grep -E "(CPU|Memory)" | head -n 10
[    0.000000] CPU: Intel(R) Core(TM) i5-8300H CPU @ 2.30GHz
[    0.000000] CPU type: Intel CoffeeLake processor
[    0.000000] CPU stepping: 10
[    0.000000] *****
[    0.000000] Hardware Thread Topology
[    0.000000] *****
[    0.000000] Sockets: 1
[    0.000000] Cores per socket: 4
[    0.000000] Threads per core: 2
[    0.000000] -----
[    0.000000] HWThread      Thread      Core      Socket      Available
[    0.000000] 0              0          0          0            *
[    0.000000] 1              0          1          0            *
[    0.000000] 2              0          2          0            *
[    0.000000] 3              0          3          0            *
[    0.000000] 4              1          0          0            *
[    0.000000] 5              1          1          0            *
[    0.000000] 6              1          2          0            *
[    0.000000] 7              1          3          0            *
[    0.000000] -----
[    0.000000] Socket 0:      ( 0 4 1 5 2 6 3 7 )
[    0.000000] -----
[    0.000000] Cache Topology
[    0.000000] *****
[    0.000000] Level:         1
[    0.000000] Size:          32 kB
[    0.000000] Cache groups:  ( 0 4 ) ( 1 5 ) ( 2 6 ) ( 3 7 )
[    0.000000] -----
[    0.000000] Level:         2
[    0.000000] Size:          256 kB
[    0.000000] Cache groups:  ( 0 4 ) ( 1 5 ) ( 2 6 ) ( 3 7 )
[    0.000000] -----
[    0.000000] Level:         3
[    0.000000] Size:          8 MB
[    0.000000] Cache groups:  ( 0 4 1 5 2 6 3 7 )
[    0.000000] -----
[    0.000000] NUMA Topology
[    0.000000] *****
[    0.000000] NUMA domains:  1
[    0.000000] -----
[    0.000000] Domain:        0
[    0.000000] Processors:    ( 0 1 2 3 4 5 6 7 )
[    0.000000] Distances:     10
[    0.000000] Free memory:   3546.2 MB
[    0.000000] Total memory:  7831.84 MB
[    0.000000] Local memory:  3915.92 MB

```

```

*****
Graphical Topology
*****
Socket 0:
+-----+
| +-----+ | +-----+ | +-----+ | +-----+ | | | | | | | | |
| | 0 4 | | | 1 5 | | | 2 6 | | | 3 7 | |
| +-----+ | +-----+ | +-----+ | +-----+ |
| +-----+ | +-----+ | +-----+ | +-----+ |
| | 32 kB | | | 32 kB | | | 32 kB | | | 32 kB | |
| +-----+ | +-----+ | +-----+ | +-----+ |
| +-----+ | +-----+ | +-----+ | +-----+ |
| | 256 kB | | | 256 kB | | | 256 kB | | | 256 kB | |
| +-----+ | +-----+ | +-----+ | +-----+ |
| +-----+ | +-----+ | +-----+ | +-----+ |
| | 8 MB | | | | | | | |
| +-----+ | +-----+ | +-----+ | +-----+ |
+-----+

```

No of nodes: 12 (4 for each as written in the machine file).

Serial Code:

```
#include "mpi.h"
#include <stdio.h>
#include <stdlib.h>
#include <float.h>
int main(int argc, char *argv[])
{
    double start,end;
    MPI_Init(&argc,&argv);
    start=MPI_Wtime();
    int i = 1, n = 100000;
    double a[n], sum = 0;
    for(i=0;i<n;i++)
    {
        a[i] = (i+1)*5.658;
        sum += a[i]
    }
    end=MPI_Wtime();
    printf("\nTime= %f",end-start);
    return 0;
}
```

Parallel Code : [MPI_Reduce]

```
#include <mpi.h>
#include <stdio.h>
#include <math.h>
#include <stdlib.h>

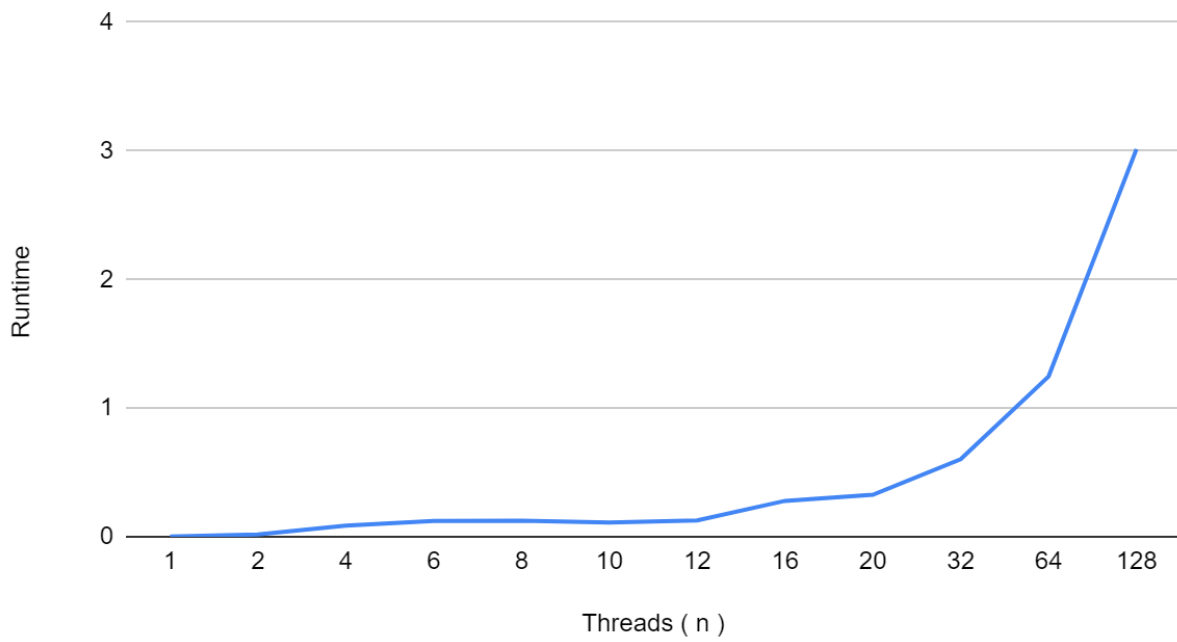
#define ARRAY_SIZE 100000

int main(int argc, char *argv[])
{
    int myid, numprocs, s;
    double vector1[ARRAY_SIZE + 50], sum = 0, vector1Receive[ARRAY_SIZE + 50],
    fSum =
    0;
    MPI_Init(&argc, &argv);
    MPI_Comm_size(MPI_COMM_WORLD, &numprocs);
    MPI_Comm_rank(MPI_COMM_WORLD, &myid);
    double start = MPI_Wtime();
    for (int i = 0; i < ARRAY_SIZE; i++)
        vector1[i] = i + 1;
    if (myid == 0)
        s = (int)floor(ARRAY_SIZE / numprocs);
    MPI_Bcast(&s, 1, MPI_INT, 0, MPI_COMM_WORLD);
    MPI_Scatter(vector1, s, MPI_DOUBLE, vector1Receive, s, MPI_DOUBLE, 0,
    MPI_COMM_WORLD);
    for (int i = 0; i < s; i++)
        sum += vector1Receive[i];
    MPI_Reduce(&sum, &fSum, 1, MPI_DOUBLE, MPI_SUM, 0, MPI_COMM_WORLD);
    if (myid == 0)
    {
        for (int i = s * numprocs; i < ARRAY_SIZE; i++)
            fSum += vector1[i];
        printf("%f\n", fSum);
        double end = MPI_Wtime();
        printf("%f\n", end - start);
    }
    MPI_Finalize();
}
```

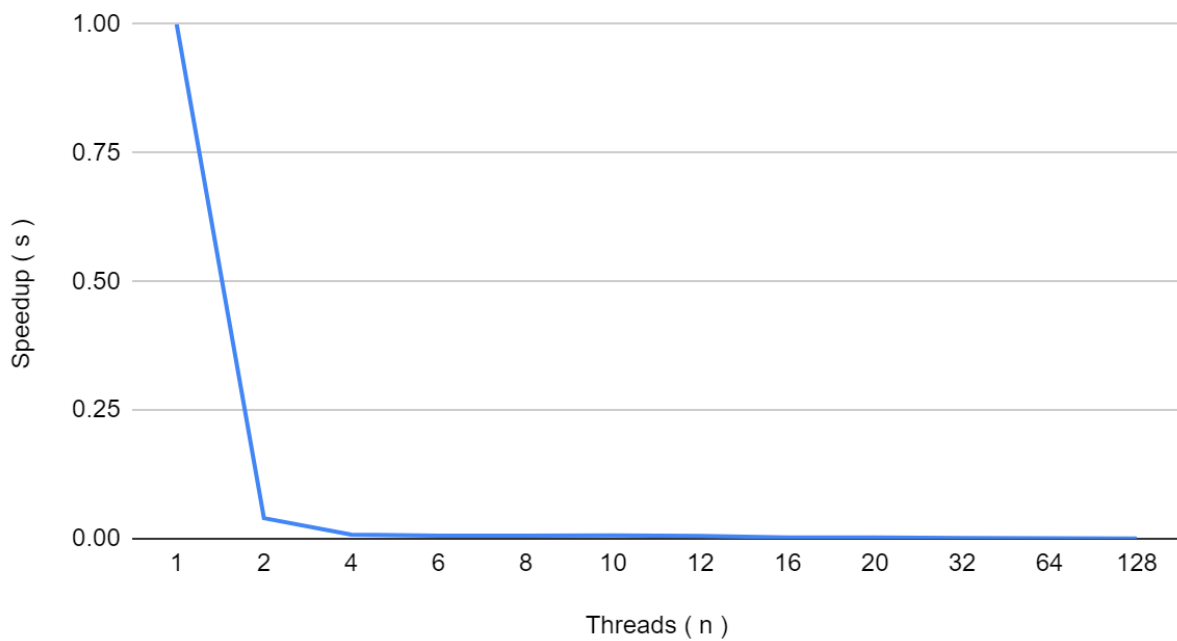
Observations : [MPI_Reduce]

Processes (n)	Runtime	Speedup (s)
1	0.000698	1
2	0.017489	0.03991080107
4	0.086921	0.00803028037
6	0.121665	0.005737064891
8	0.124371	0.005612240796
10	0.110482	0.006317771221
12	0.126299	0.005526567906
16	0.278594	0.002505438021
20	0.325962	0.0021413539
32	0.602774	0.001157979608
64	1.246792	0.0005598367651
128	3.013827	0.0002315992258

Runtime vs. Processes(n)



Speedup (s) vs. Processes(n)



Parallel Code : [Without MPI_Reduce]

```
#include <mpi.h>
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>

// size of array
#define n 10000

int a[10000];

// Temporary array for slave process
int a2[1000000000];

int main(int argc, char* argv[])
{
    double start,end;
    int pid, np,
        elements_per_process,
        n_elements_recieved;
    // np -> no. of processes
    // pid -> process id

    MPI_Status status;

    // Creation of parallel processes
    MPI_Init(&argc, &argv);
    start=MPI_Wtime();
    for (int i = 0; i < 10000; i++)
        a[i] = (i + 1)*8.754;
    // find out process ID,
    // and how many processes were started
    MPI_Comm_rank(MPI_COMM_WORLD, &pid);
    MPI_Comm_size(MPI_COMM_WORLD, &np);

    // master process
    if (pid == 0) {
        int index, i;
        elements_per_process = n / np;
```

```
// check if more than 1 processes are run
if (np > 1) {
    // distributes the portion of array
    // to child processes to calculate
    // their partial sums
    for (i = 1; i < np - 1; i++) {
        index = i * elements_per_process;

        MPI_Send(&elements_per_process,
                 1, MPI_INT, i, 0,
                 MPI_COMM_WORLD);
        MPI_Send(&a[index],
                 elements_per_process,
                 MPI_INT, i, 0,
                 MPI_COMM_WORLD);
    }

    // last process adds remaining elements
    index = i * elements_per_process;
    int elements_left = n - index;

    MPI_Send(&elements_left,
             1, MPI_INT,
             i, 0,
             MPI_COMM_WORLD);
    MPI_Send(&a[index],
             elements_left,
             MPI_INT, i, 0,
             MPI_COMM_WORLD);
}

// master process add its own sub array
int sum = 0;
for (i = 0; i < elements_per_process; i++)
    sum += a[i];

// collects partial sums from other processes
int tmp;
for (i = 1; i < np; i++) {
    MPI_Recv(&tmp, 1, MPI_INT,
```

```
        MPI_ANY_SOURCE, 0,
        MPI_COMM_WORLD,
        &status);
    int sender = status.MPI_SOURCE;

    sum += tmp;
}

// prints the final sum of array
printf("Sum of array is : %d\n", sum);
end=MPI_Wtime();
printf("\nTime_hello= %f",end-start);
}
// slave processes
else {
    MPI_Recv(&n_elements_recieved,
            1, MPI_INT, 0, 0,
            MPI_COMM_WORLD,
            &status);

    // stores the received array segment
    // in local array a2
    MPI_Recv(&a2, n_elements_recieved,
            MPI_INT, 0, 0,
            MPI_COMM_WORLD,
            &status);

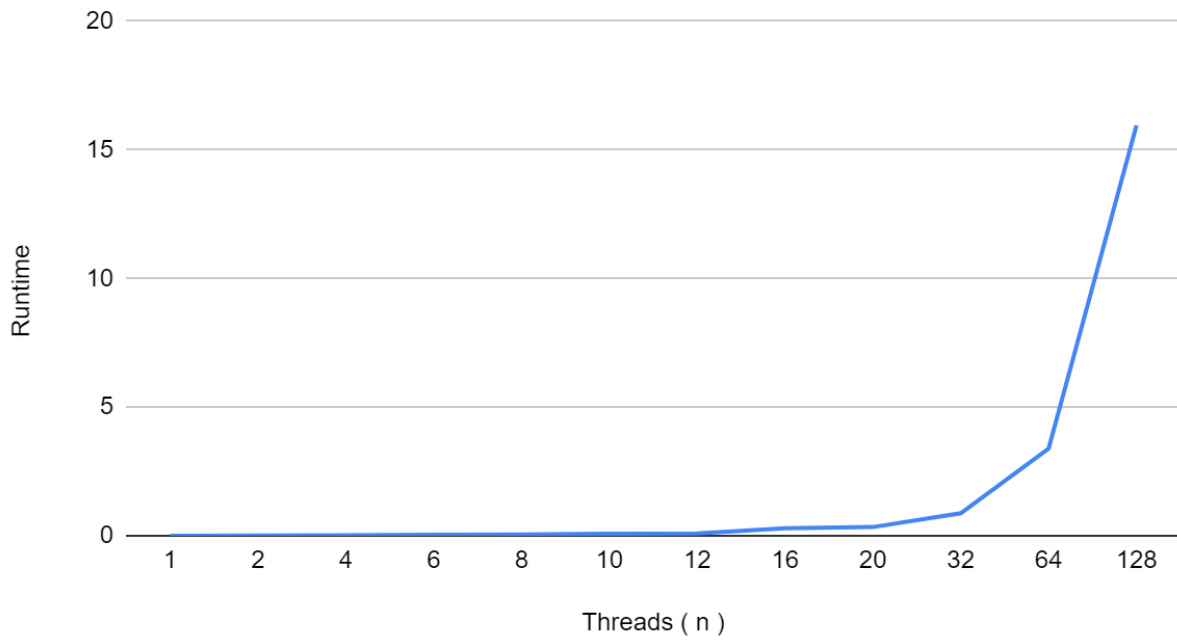
    // calculates its partial sum
    int partial_sum = 0;
    for (int i = 0; i < n_elements_recieved; i++)
        partial_sum += a2[i];

    // sends the partial sum to the root process
    MPI_Send(&partial_sum, 1, MPI_INT,
            0, 0, MPI_COMM_WORLD);
}
// cleans up all MPI state before exit of process
MPI_Finalize();
return 0;
}
```

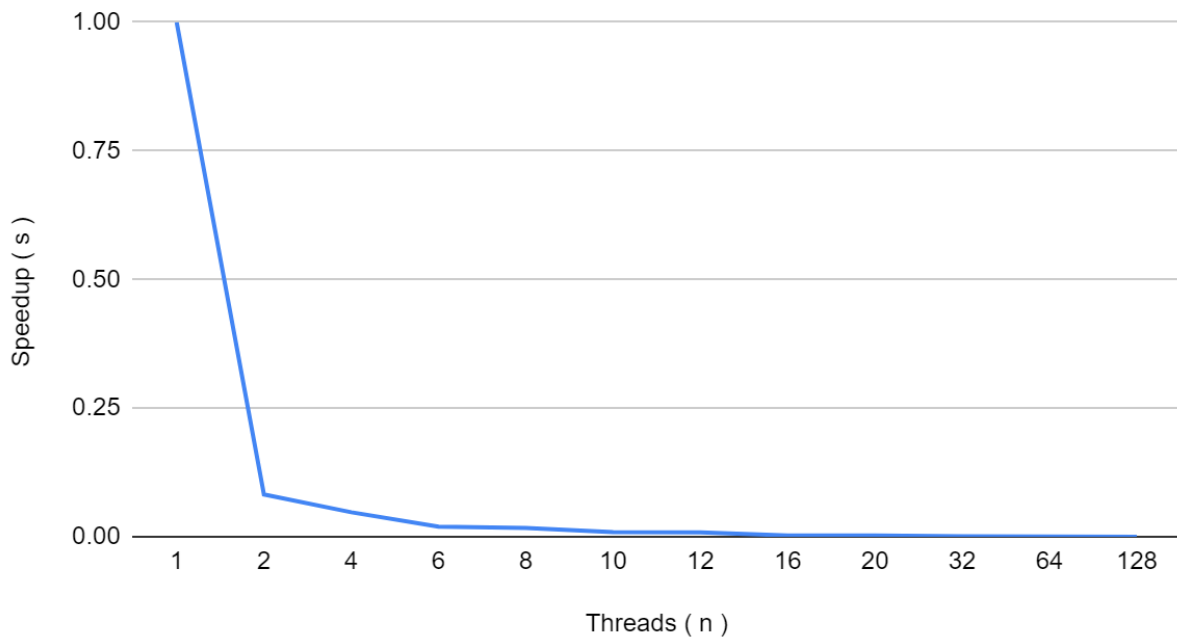

Observations : [Without MPI_Reduce]

Processes (n)	Runtime	Speedup (s)
1	0.000698	1
2	0.008503	0.08208867459
4	0.014682	0.04754120692
6	0.036009	0.01938404288
8	0.04158	0.01678691679
10	0.078148	0.008931770487
12	0.082783	0.008431682833
16	0.290735	0.002400811736
20	0.334831	0.002084633741
32	0.873777	0.000798830823
64	3.385408	0.0002061789894
128	15.950274	0.00004376100373

Runtime vs. Processes (n)



Speedup (s) vs. Processes (n)



Inference: (**Note:** Execution time, graph, and inference will be based on hardware configuration)

- Since MPI is a distributed memory architecture, the communication overhead between nodes causes the parallel code to run slower compared to serial code (running in 1 node or only in master).