

HPC LAB - Matrix Addition

Name: Paleti Krishnasai

Roll No: CED18I039

Programming Environment: OpenMP

Problem: Matrix Addition

Date: 19th August 2021

Hardware Configuration:

PU NAME: Intel(R) Core(TM) i5-8300H CPU @ 2.30GHz

Number of Sockets: 1

Cores per Socket : 4

Threads per core: 2

L1d cache: 128 KiB

L1i cache: 128 KiB

L2 cache: 1 MiB


L3 cache: 8 MiB

CPU-Z - ID : hzkq5i

CPU | Caches | Mainboard | Memory | SPD | Graphics | Bench | About

Processor

Name	Intel Core i5 8300H		
Code Name	Coffee Lake	Max TDP	45 W
Package	Socket 1440 FCBGA		
Technology	14 nm	Core Voltage	1.08 V



Specification Intel(R) Core(TM) i5-8300H CPU @ 2.30GHz

Family	6	Model	E	Stepping	A
Ext. Family	6	Ext. Model	9E	Revision	U0

Instructions MMX, SSE, SSE2, SSE3, SSE4.1, SSE4.2, EM64T, VT-x, AES, AVX, AVX2, FMA3

Clocks (Core #0)

Core Speed	3886.66 MHz
Multiplier	x 39.0 (8 - 40)
Bus Speed	99.66 MHz
Rated FSB	

Cache

L1 Data	4 x 32 KBytes	8-way
L1 Inst.	4 x 32 KBytes	8-way
Level 2	4 x 256 KBytes	4-way
Level 3	8192 KBytes	16-way

Selection Processor #1 Cores 4 Threads 8

CPU-Z Ver. 1.94.0.x64 Tools Validate Close

```
paleti@paleti-Lenovo-ideapad-330-15ICH:~$ lscpu
Architecture:          x86_64
CPU op-mode(s):        32-bit, 64-bit
Byte Order:            Little Endian
Address sizes:          39 bits physical, 48 bits virtual
CPU(s):                8
On-line CPU(s) list:    0-7
Thread(s) per core:     2
Core(s) per socket:     4
Socket(s):              1
NUMA node(s):           1
Vendor ID:              GenuineIntel
CPU family:             6
Model:                  158
Model name:             Intel(R) Core(TM) i5-8300H CPU @ 2.30GHz
Stepping:               10
CPU MHz:                900.021
CPU max MHz:            4000.0000
CPU min MHz:            800.0000
BogoMIPS:               4599.93
Virtualization:         VT-x
L1d cache:              128 KiB
L1i cache:              128 KiB
L2 cache:               1 MiB
L3 cache:               8 MiB
NUMA node0 CPU(s):      0-7
Vulnerability Itlb multihit: KVM: Mitigation: VMX disabled
Vulnerability L1tf:       Mitigation; PTE Inversion; VMX conditional cache flushes, SMT vulnerable
Vulnerability Mds:         Mitigation; Clear CPU buffers; SMT vulnerable
Vulnerability Meltdown:    Mitigation; PTI
Vulnerability Spec store bypass: Mitigation; Speculative Store Bypass disabled via prctl and seccomp
Vulnerability Spectre v1:  Mitigation; usercopy/swapgs barriers and __user pointer sanitization
Vulnerability Spectre v2:  Mitigation; Full generic retpoline, IBPB conditional, IBRS_FW, STIBP conditional, RSB filling
Vulnerability Srbds:       Mitigation; Microcode
Vulnerability Tsx async abort: Not affected
```

Serial Code:

```
#include <stdio.h>
#include <stdlib.h>
#include <omp.h>
#include <time.h>

#define size 400

int main()
{
    Double
M1[size][size]={1075.75},M2[size][size]={1594.97},M3[size][size];
    float startTime,endTime,execTime;
    int i,j;
    int omp_rank;

    startTime = omp_get_wtime();
    {
        omp_rank = omp_get_thread_num();
        for(i=0;i<400;i++)
        {
            for(j=0;j<400;j++)
            {
                for(int k=0;k<1000000;k++)
                    M3[i][j] = M1[i][j] + M2[i][j];
            }
        }
    }
    endTime = omp_get_wtime();

    execTime = endTime-startTime;
    printf("%f \n",execTime);
    return(0);
}
```

Parallel Code :

```
#include <stdio.h>
#include <stdlib.h>
#include <omp.h>
#include <time.h>

#define size 400

int main()
{
    long double
M1[size][size]={1075.75},M2[size][size]={1594.97},M3[size][size];
    float startTime,endTime,execTime;
    int i,j;
    int omp_rank;

    startTime = omp_get_wtime();
    #pragma omp parallel private (i,j) shared (M1,M2,M3)
    {
        omp_rank = omp_get_thread_num();
        #pragma omp for
        for(i=0;i<250;i++)
        {
            for(j=0;j<250;j++)
            {
                for(int k=0;k<100000;k++)
                    M3[i][j] = M1[i][j] + M2[i][j];
            }
        }
    }
    endTime = omp_get_wtime();

    execTime = endTime-startTime;
    printf("%f \n",execTime);
    return(0);
}
```

Compilation and Execution:

To enable OpenMP environment, use -fopenmp flag while compiling using gcc.

```
gcc -O0 -fopenmp matrix_add_mp.c
```

Then,

export OMP_NUM_THREADS= no of threads for parallel execution.

./a.out

Observations :

Threads (n)	Runtime	Speedup (s)	Parallelization Fraction
1	55.125	1	
2	27.640625	1.994347089	0.9971655329
4	17.144531	3.215311052	0.9186507997
8	14.648438	3.763199872	0.8391642268
16	14.230469	3.873730374	0.7913076293
32	13.699219	4.023952022	0.775729642
64	13.949219	3.951834149	0.7588093392
128	13.53125	4.073903002	0.760476369

Speed up can be found using the following formula,

$$S(n)=T(1)/T(n)$$

where, **S(n)** = Speedup for thread count 'n'

T(1) = Execution Time for Thread count '1' (serial code)

T(n) = Execution Time for Thread count 'n' (serial code)

Parallelization Fraction can be found using the following formula,

$$S(n)=1/((1 - p) + p/n)$$

where, **S(n)** = Speedup for thread count 'n'

n = Number of threads


p = Parallelization fraction

Assumption:


Following extra for loop is added to increase the number of operations in the parallel region to visualize the effect of multi-threading in matrix addition.

```
for(int k=0;k<10000;k++)
```

Runtime vs. Threads (n)



Speedup (s) vs. Threads (n)



Inference: (Note: Execution time, graph and inference will be based on hardware configuration)

- At thread count 128 maximum speedup is observed.
- The Runtime is least at thread count is 128, but it fluctuates from there on rising slightly. This may be due to the fact that there are 4 cores and 2 threads on each core are active at a moment for this task concurrently.