Paleti Krishnasai
CED18I039

HPC LAB - Matrix Multiplication

Name: Paleti Krishnasai
Roll No: CED18I039
Programming Environment: OpenMP
Problem: Matrix Multiplication
Date: 19th August 2021

Hardware Configuration:
PU NAME: Intel(R) Core(TM) i5-8300H CPU @ 2.30GHz
Number of Sockets: 1
Cores per Socket : 4
Threads per core: 2
L1d cache: 128 KiB
L1i cache: 128 KiB
L2 cache: 1 MiB
L3 cache: 8 MiB

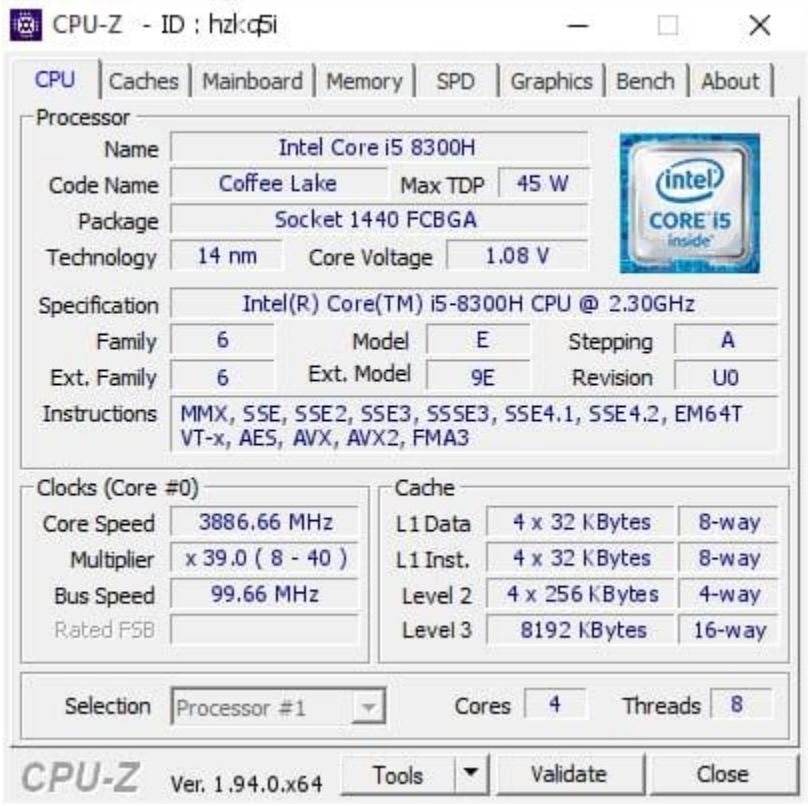

Paleti Krishnasai
CED18I039

```
paleti@paleti-Lenovo-ideapad-330-15ICH:~$ lscpu
Architecture:                    x86_64
CPU op-mode(s):                  32-bit, 64-bit
Byte Order:                      Little Endian
Address sizes:                   39 bits physical, 48 bits virtual
CPU(s):                          8
On-line CPU(s) list:             0-7
Thread(s) per core:              2
Core(s) per socket:              4
Socket(s):                       1
NUMA node(s):                    1
Vendor ID:                       GenuineIntel
CPU family:                      6
Model:                           158
Model name:                      Intel(R) Core(TM) i5-8300H CPU @ 2.30GHz
Stepping:                        10
CPU MHz:                         900.021
CPU max MHz:                     4000.0000
CPU min MHz:                     800.0000
BogoMIPS:                        4599.93
Virtualization:                  VT-x
L1d cache:                       128 KiB
L1i cache:                       128 KiB
L2 cache:                        1 MiB
L3 cache:                        8 MiB
NUMA node0 CPU(s):               0-7
Vulnerability Itlb multihit:     KVM: Mitigation: VMX disabled
Vulnerability L1tf:              Mitigation; PTE Inversion; VMX conditional cach
                                 e flushes, SMT vulnerable
Vulnerability Mds:               Mitigation; Clear CPU buffers; SMT vulnerable
Vulnerability Meltdown:          Mitigation; PTI
Vulnerability Spec store bypass: Mitigation; Speculative Store Bypass disabled v
                                 ia prctl and seccomp
Vulnerability Spectre v1:        Mitigation; usercopy/swapgs barriers and __user
                                  pointer sanitization
Vulnerability Spectre v2:        Mitigation; Full generic retpoline, IBPB condit
                                 ional, IBRS_FW, STIBP conditional, RSB filling
Vulnerability Srbds:             Mitigation; Microcode
Vulnerability Tsx async abort:   Not affected
```

Paleti Krishnasai
CED18I039

Serial Code:

```
//row major matrix multiplication
#include <stdio.h>
#include <stdlib.h>
#include <omp.h>
#include <time.h>

#define n 350

int main()
{
   int i,j;
   double M1[n][n]={984.46},M2[n][n]={634.52},M3[n][n];
   float startTime, endTime, execTime;
   int omp_rank;

   startTime = omp_get_wtime();
   {
       for(i=0 ; i< n ; i++)
       {
           omp_rank = omp_get_thread_num();
           for(j=0 ; j< n ;j++)
           {
               M1[i][j] = 0;
               for(int k=0;k< n; k++)
               {
                   for(int k=0;k<100;k++)
                   M3[i][j] = M3[i][j] + M1[i][k] * M2[k][j];
               }
           }
       }
   }

   endTime = omp_get_wtime();
   execTime = endTime - startTime;
   printf("%f \n",execTime);
   return(0);
}
```

Paleti Krishnasai
CED18I039

Parallel Code :

```
//row major matrix multiplication
#include <stdio.h>
#include <stdlib.h>
#include <omp.h>
#include <time.h>

#define size 350

int main()
{
   int i,j;
   double M1[size][size]={984.46},M2[size][size]={634.52},M3[size][size];
   float startTime, endTime, execTime;
   int omp_rank;

   startTime = omp_get_wtime();
   #pragma omp parallel private (i,j) shared (M1,M2,M3)
   {
       #pragma omp for
       for(i=0 ; i< 350 ; i++)
       {
           omp_rank = omp_get_thread_num();
           for(j=0 ; j< 350 ;j++)
           {
               M1[i][j] = 0;
               for(int k=0;k< 350; k++)
               {
                   for(int k=0;k<100;k++)
                   M3[i][j] = M3[i][j] + M1[i][k] * M2[k][j];
               }
           }
       }
   }

   endTime = omp_get_wtime();
   execTime = endTime - startTime;
   printf("%f \n",execTime);
   return(0);
}
```

Paleti Krishnasai
CED18I039

Compilation and Execution:

To enable OpenMP environment, use -fopenmp flag while compiling using gcc.

gcc -O0 -fopenmp matrix_mul_mp.c

Then,

export OMP_NUM_THREADS= no of threads for parallel execution.

./a.out

Observations :

Threads (n)	Runtime	Speedup (s)	Parallelization Fraction
1	28.523438	1	
2	14.296875	1.995082002	0.9975349395
4	7.6875	3.710365919	0.9739797379
6	7.605469	3.75038515	0.8800328628
8	7.320312	3.896478456	0.849552007
16	7.699219	3.704718362	0.7787455449
32	7.945312	3.58997079	0.7447186597
64	7.359375	3.875796246	0.7537660961
128	7.308594	3.90272575	0.7496252736

Speed up can be found using the following formula,
S(n)=T(1)/T(n)
where, **S(n)** = Speedup for thread count ‘n’
T(1) = Execution Time for Thread count ‘1’ (serial code)
T(n) = Execution Time for Thread count ‘n’ (serial code)

Parallelization Fraction can be found using the following formula,
S(n)=1/((1 - p) + p/n)
where, **S(n)** = Speedup for thread count ‘n’
n = Number of threads
p = Parallelization fraction

Paleti Krishnasai
CED18I039

Assumption:

Following extra for loop is added to increase the number of operations in the parallel region to visualize the effect of multi-threading in matrix multiplication.

for(int k=0;k<100;k++)

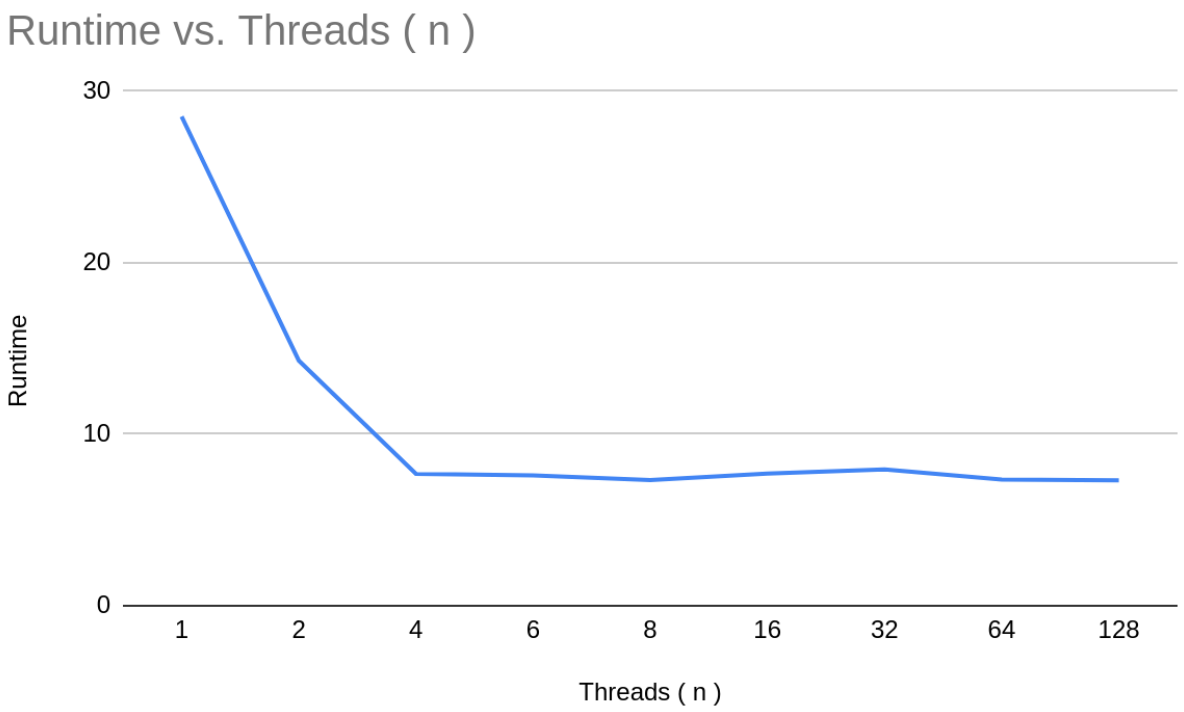

Paleti Krishnasai
CED18I039

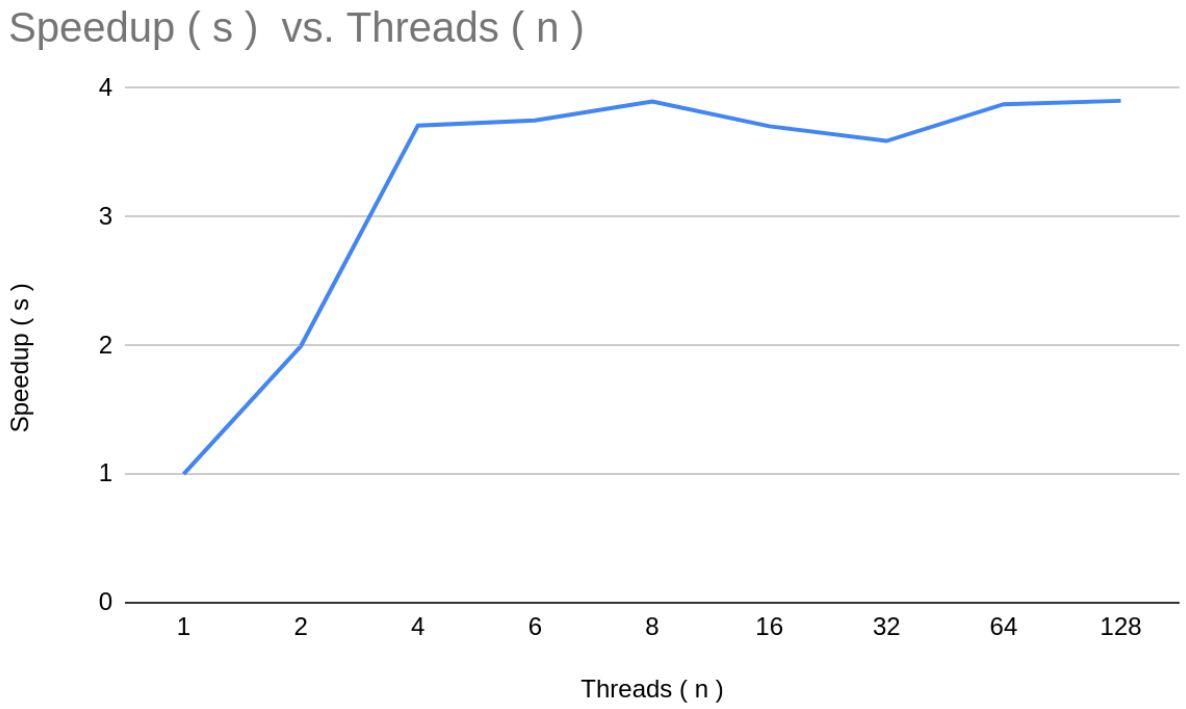

Inference: (**Note**: Execution time, graph and inference will be based on hardware configuration)

- At thread count 8, 128 maximum speedup is observed.
- The Runtime is least at thread count is 8, 128, but it fluctuates from there on rising slightly. This may be due to the fact that there are 4 cores and 2 threads on each core are active at a moment for this task concurrently.